



---

# Production-ready Flask & Django apps on Kubernetes

Jamie Hewland

PyConZA 2019

**aruba**  
a Hewlett Packard  
Enterprise company

# Hi, I'm Jamie



@jayhewland

jamie.hewland@hpe.com

Cape Town, South Africa



- ~3 years as a Site Reliability Engineer (SRE) working on container things
- Now a Backend Engineer at Aruba working on Python services
- Spoke at PyConZA 2017 about containerizing Django web apps: <https://youtu.be/T2hooQzvurQ>

- This is not an intro to Kubernetes or Python web apps
- But if you are familiar with one, this talk may help you understand the other
- There will be more YAML than Python 😭

---

About this talk

All text is in Arial

Kubernetes resources are always in a monospace font, bold, and Title Case:

- Deployment
- Pod
- ConfigMap



## AGENDA

---

### INTRO

- Why Kubernetes?
- “The WSGI Stack”

### ADAPTING TO KUBERNETES

- Handling requests
- Static file serving
- Deployments
- Configuration
- Credentials
- Database migrations

### THE FUTURE



# Why Kubernetes?

## Abstraction

- Compute/networking/storage abstracted over multiple hosts
- Developers don't think about sysadmin tasks
- Multi-cloud support
- Declarative infrastructure requirements for apps
- Improve resource usage

## Automation

- Deployments with minimal manual intervention
- Fail-over when an app misbehaves
- Tie together multiple tools

## Ecosystem

- Huge ecosystem of tools through the CNCF to extend Kubernetes
- Spend less time implementing something that has been done before



***Kubernetes is not  
the solution for  
every problem***

DevOps Thought Liker  
@dexhorthy

Follow ▾

Deployed my blog on Kubernetes

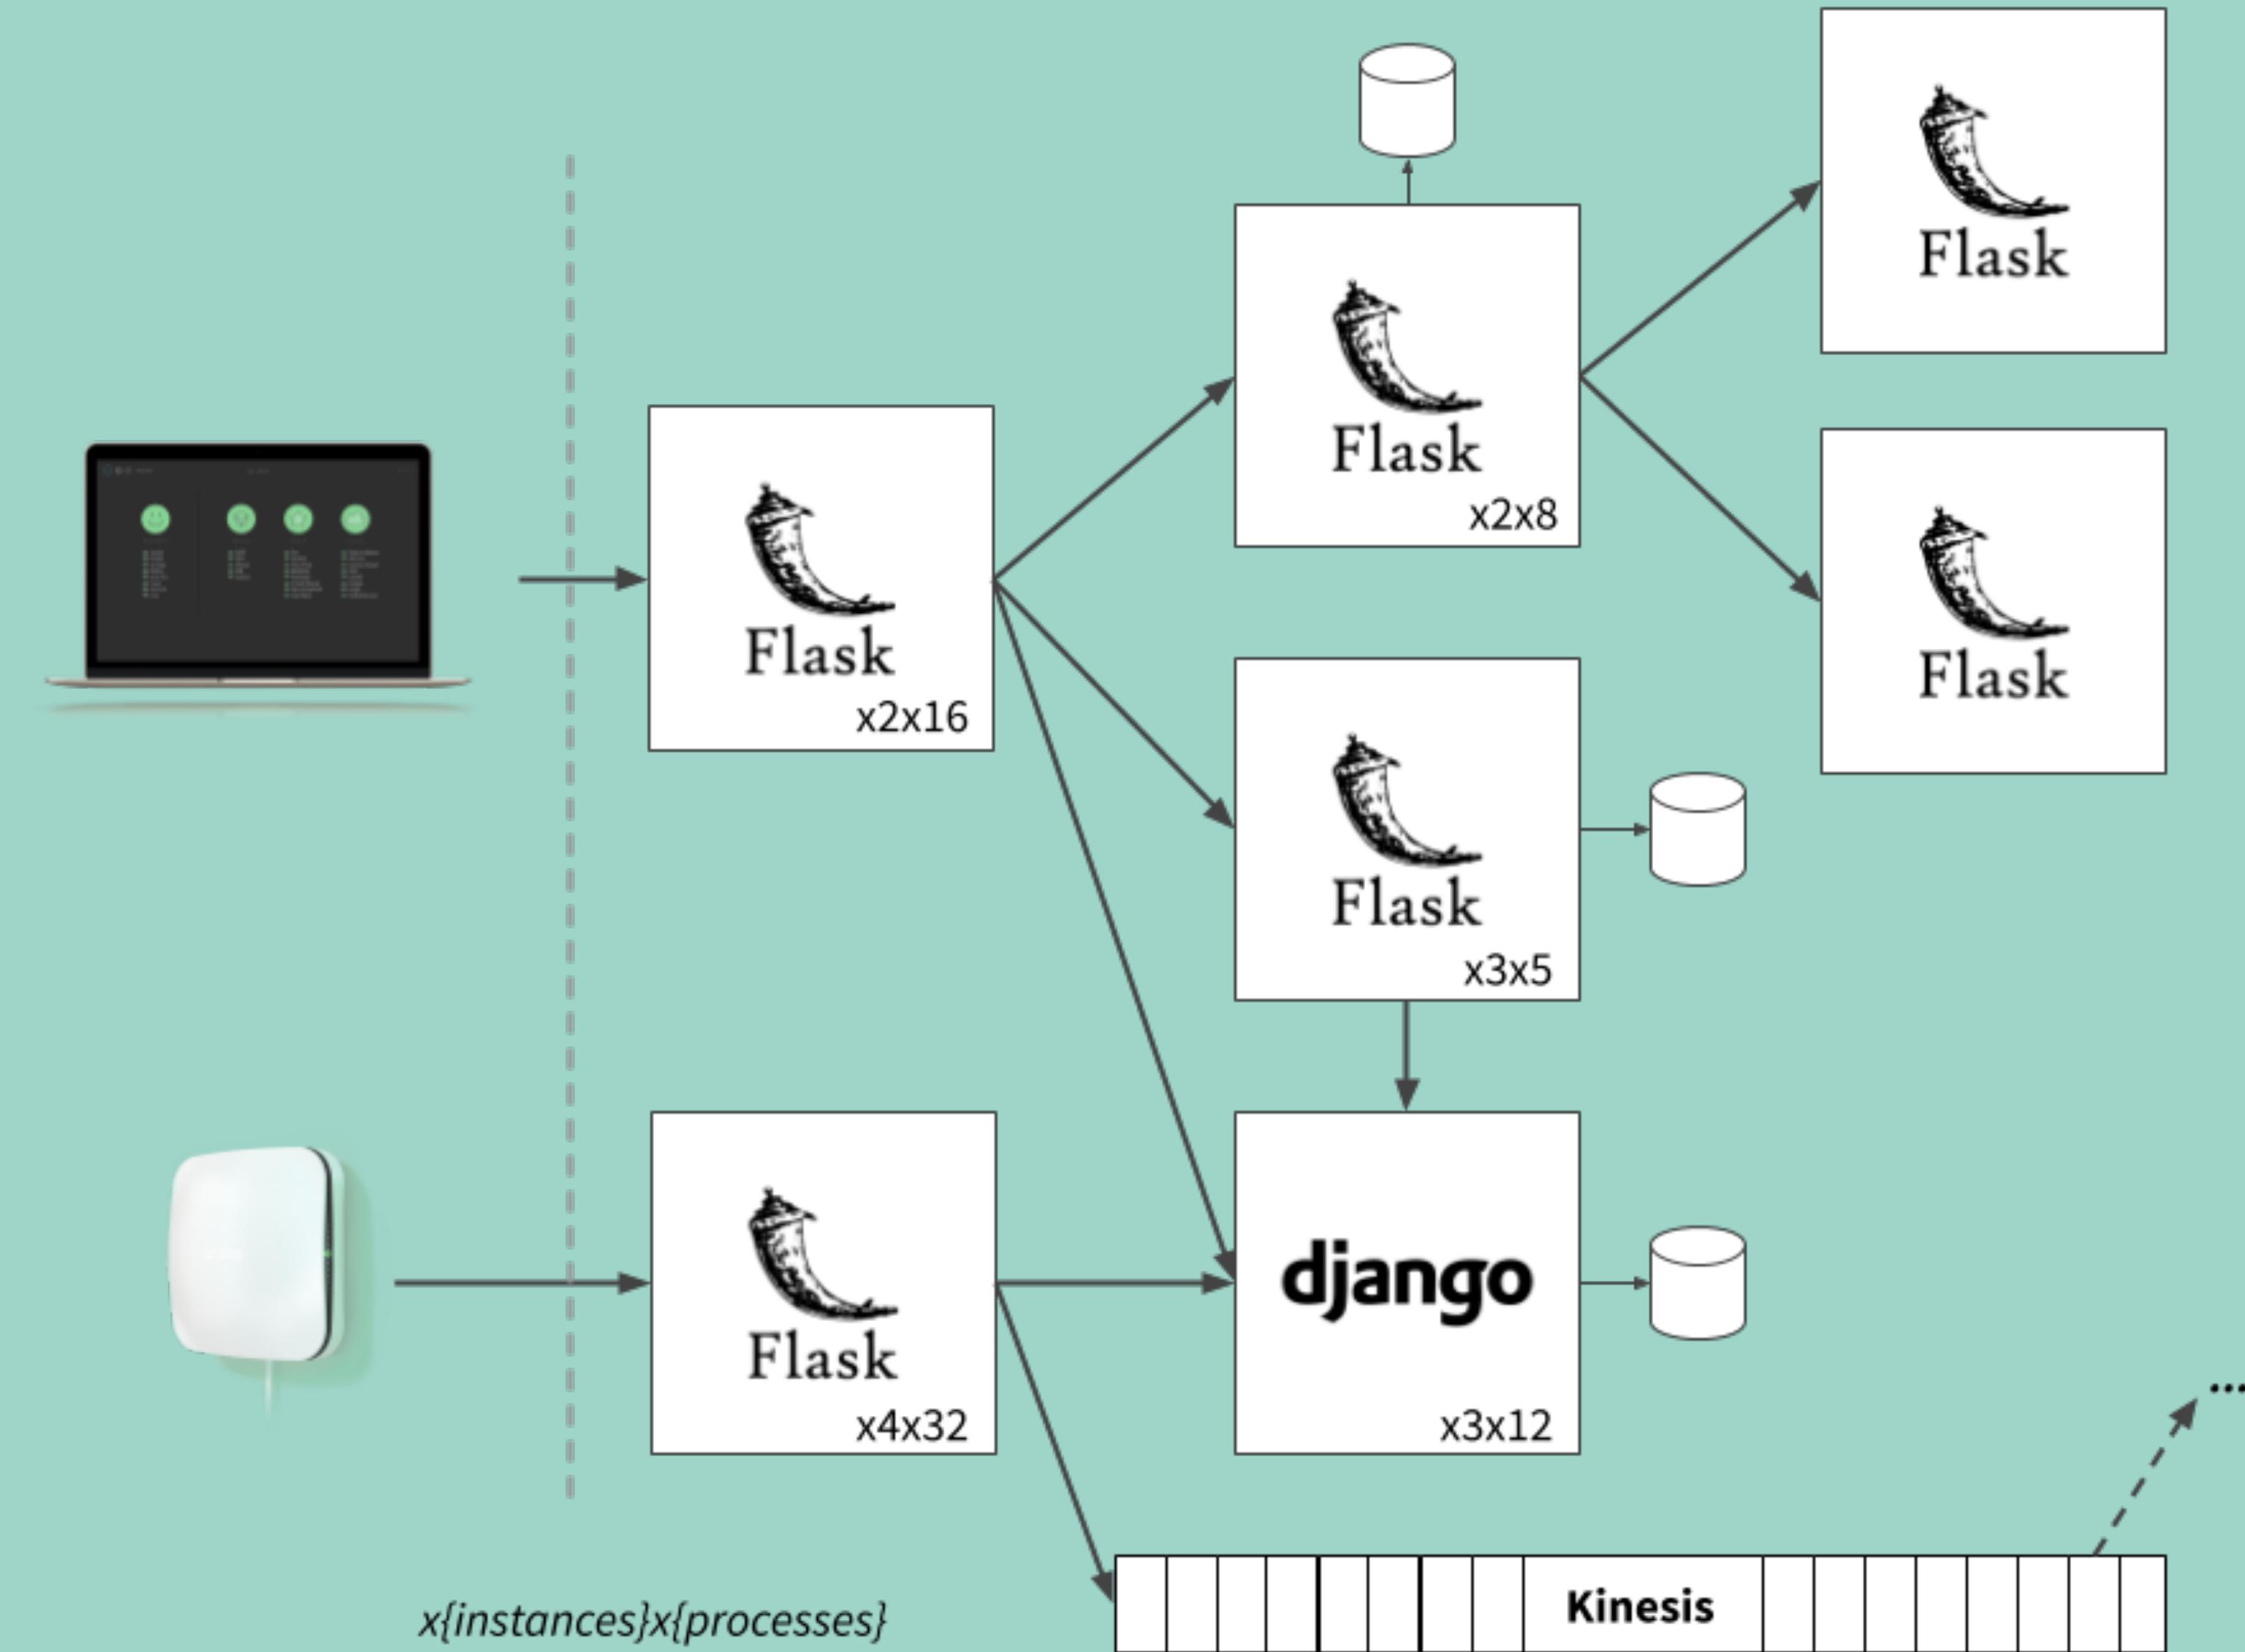


3:40 PM - 24 Apr 2017

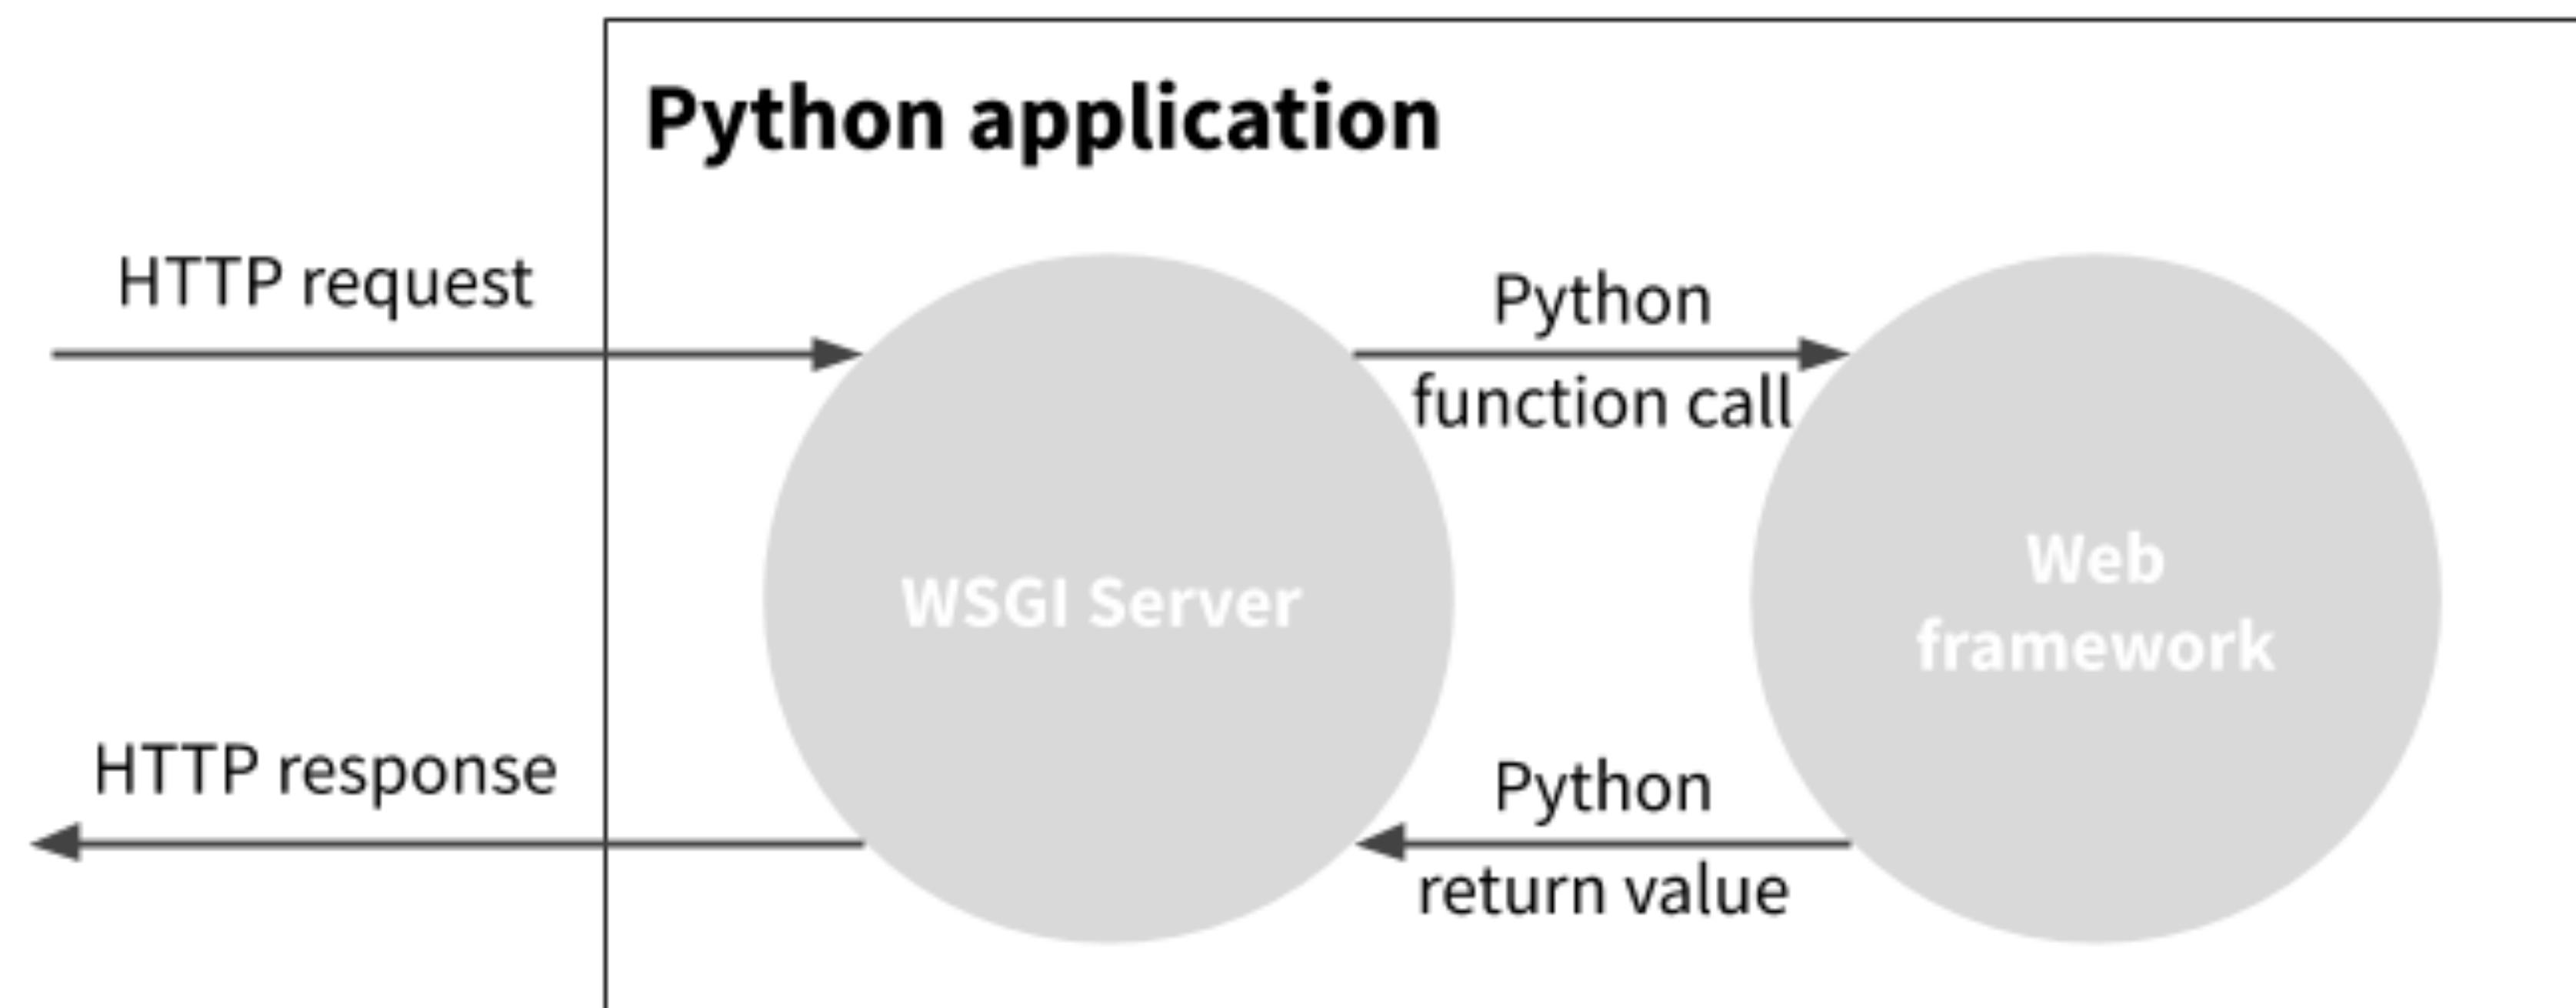
2,941 Retweets 5,853 Likes

37 2.9K 5.9K

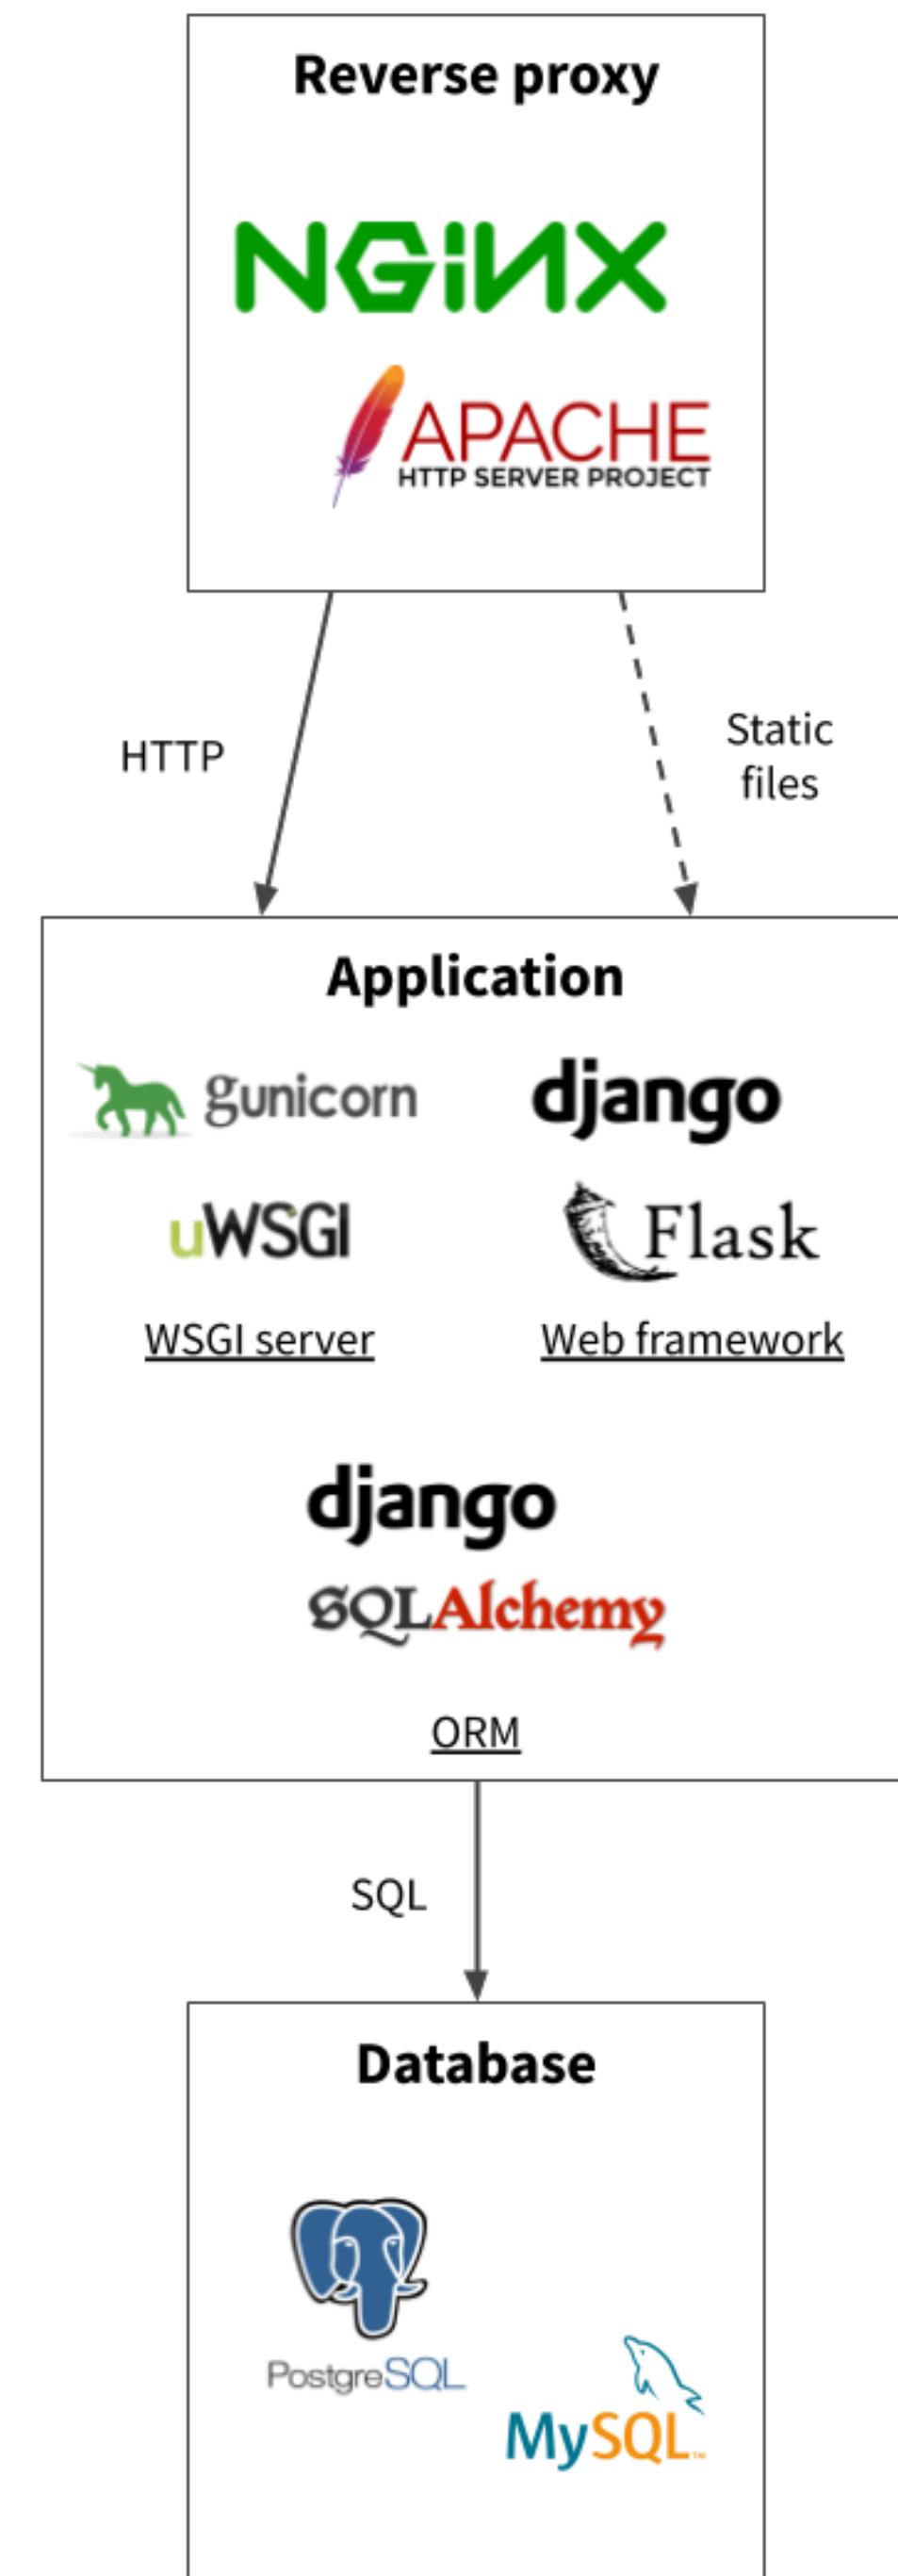
# Aruba UXI architecture



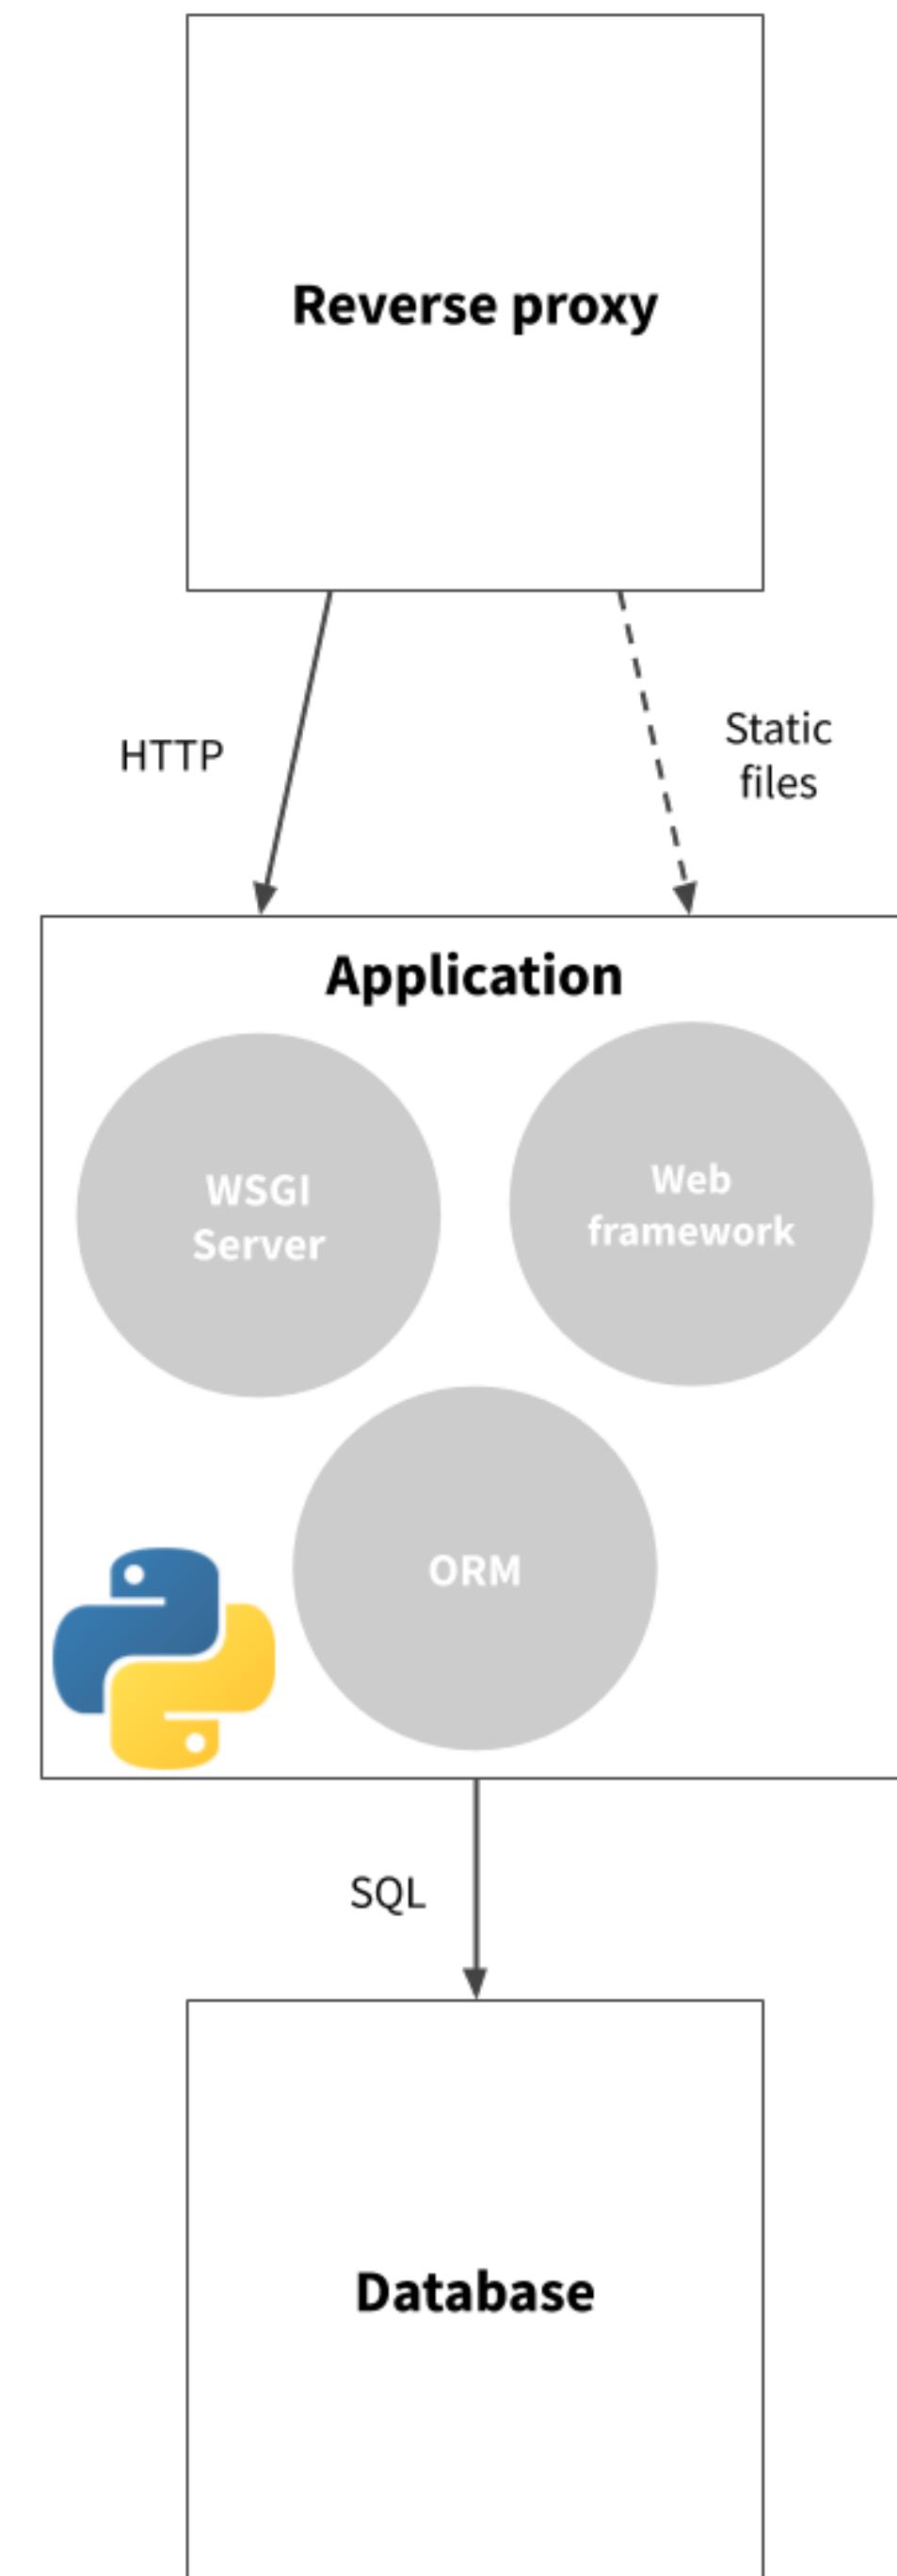
# WSGI: Web Server Gateway Interface



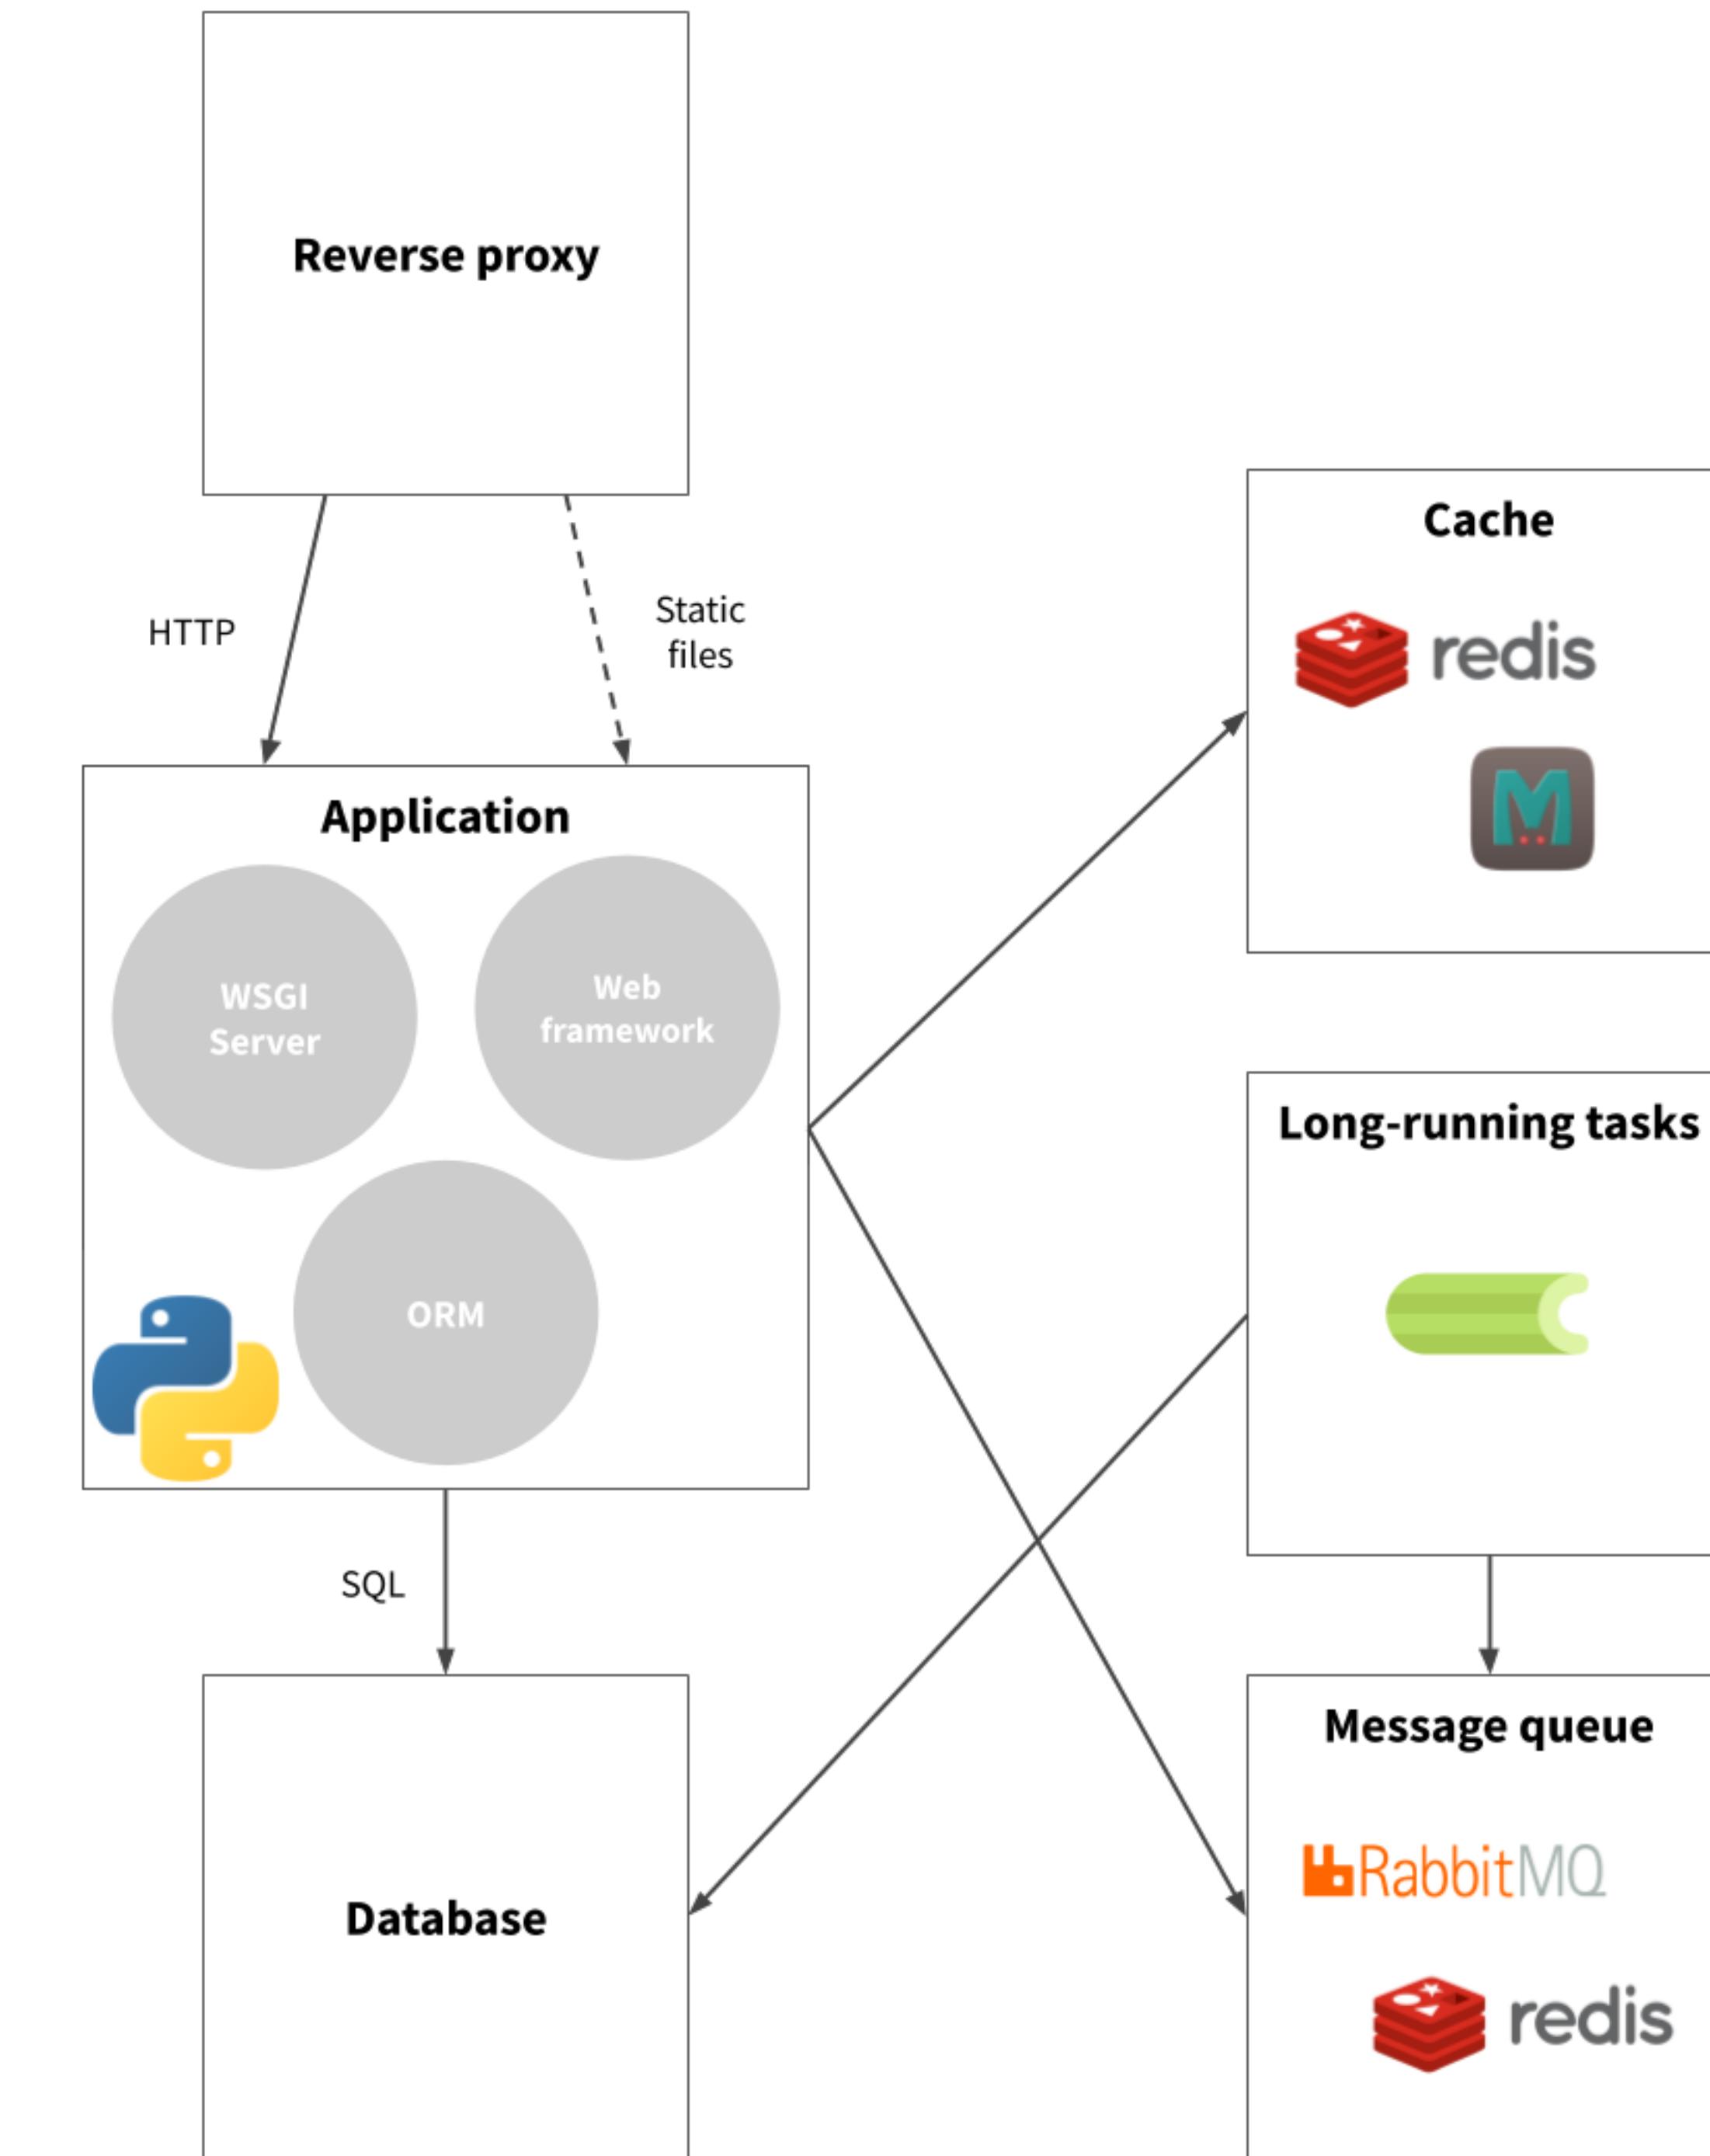
# The WSGI stack



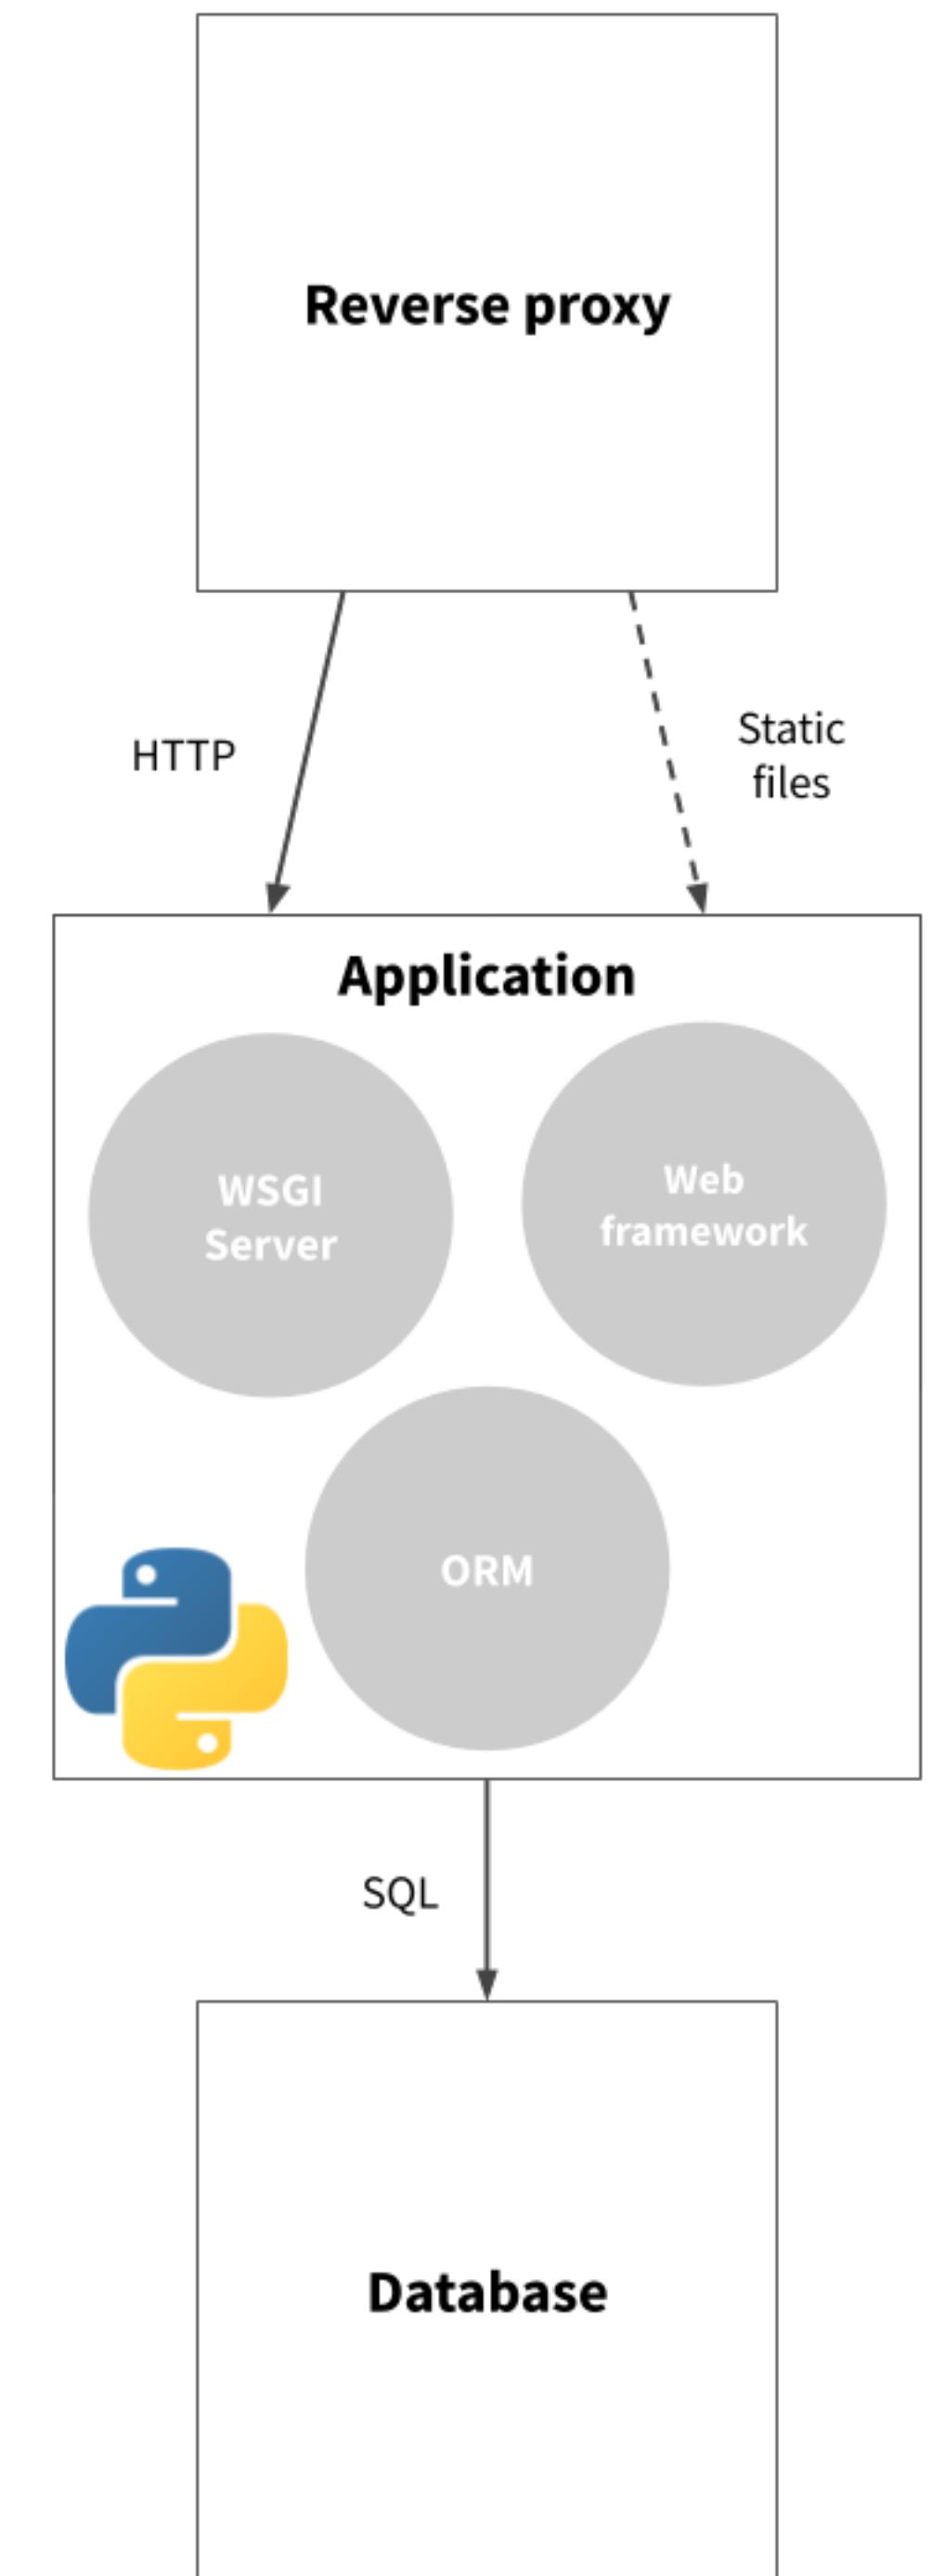
# The WSGI stack



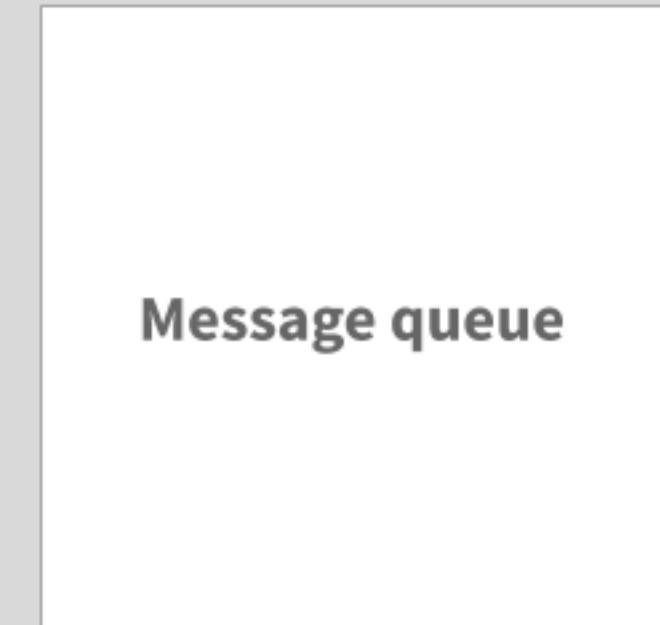
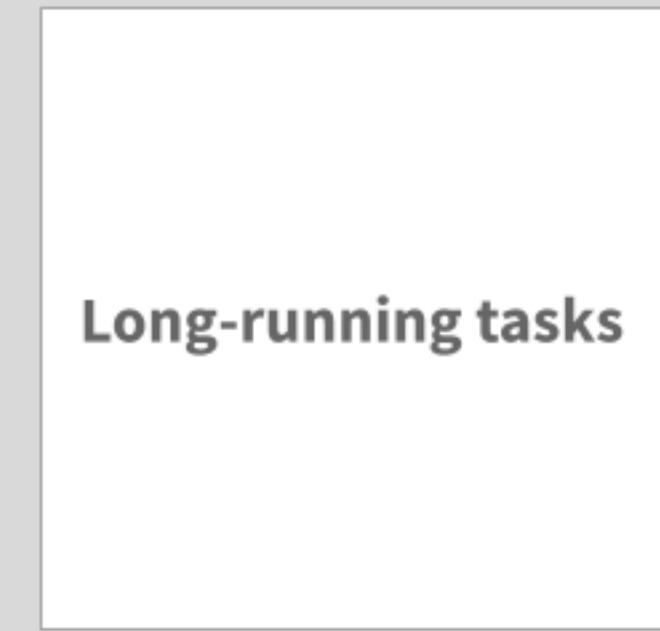
# The WSGI stack



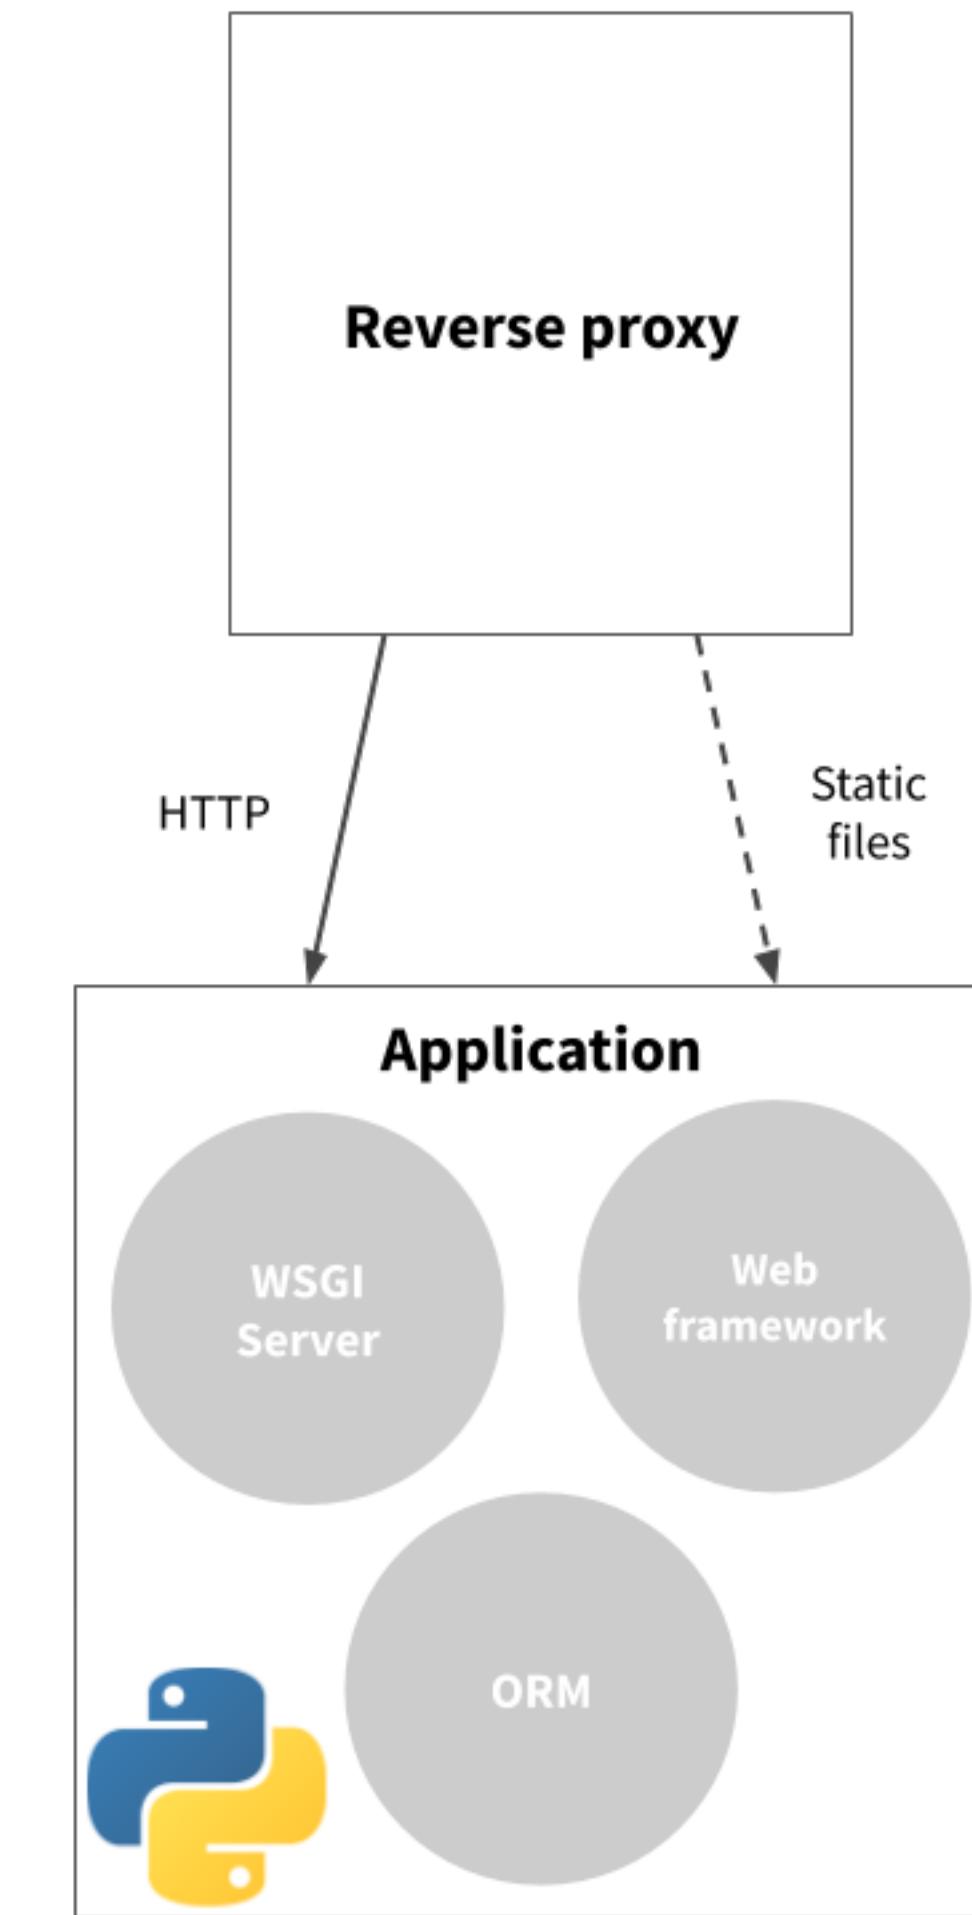
# The WSGI stack



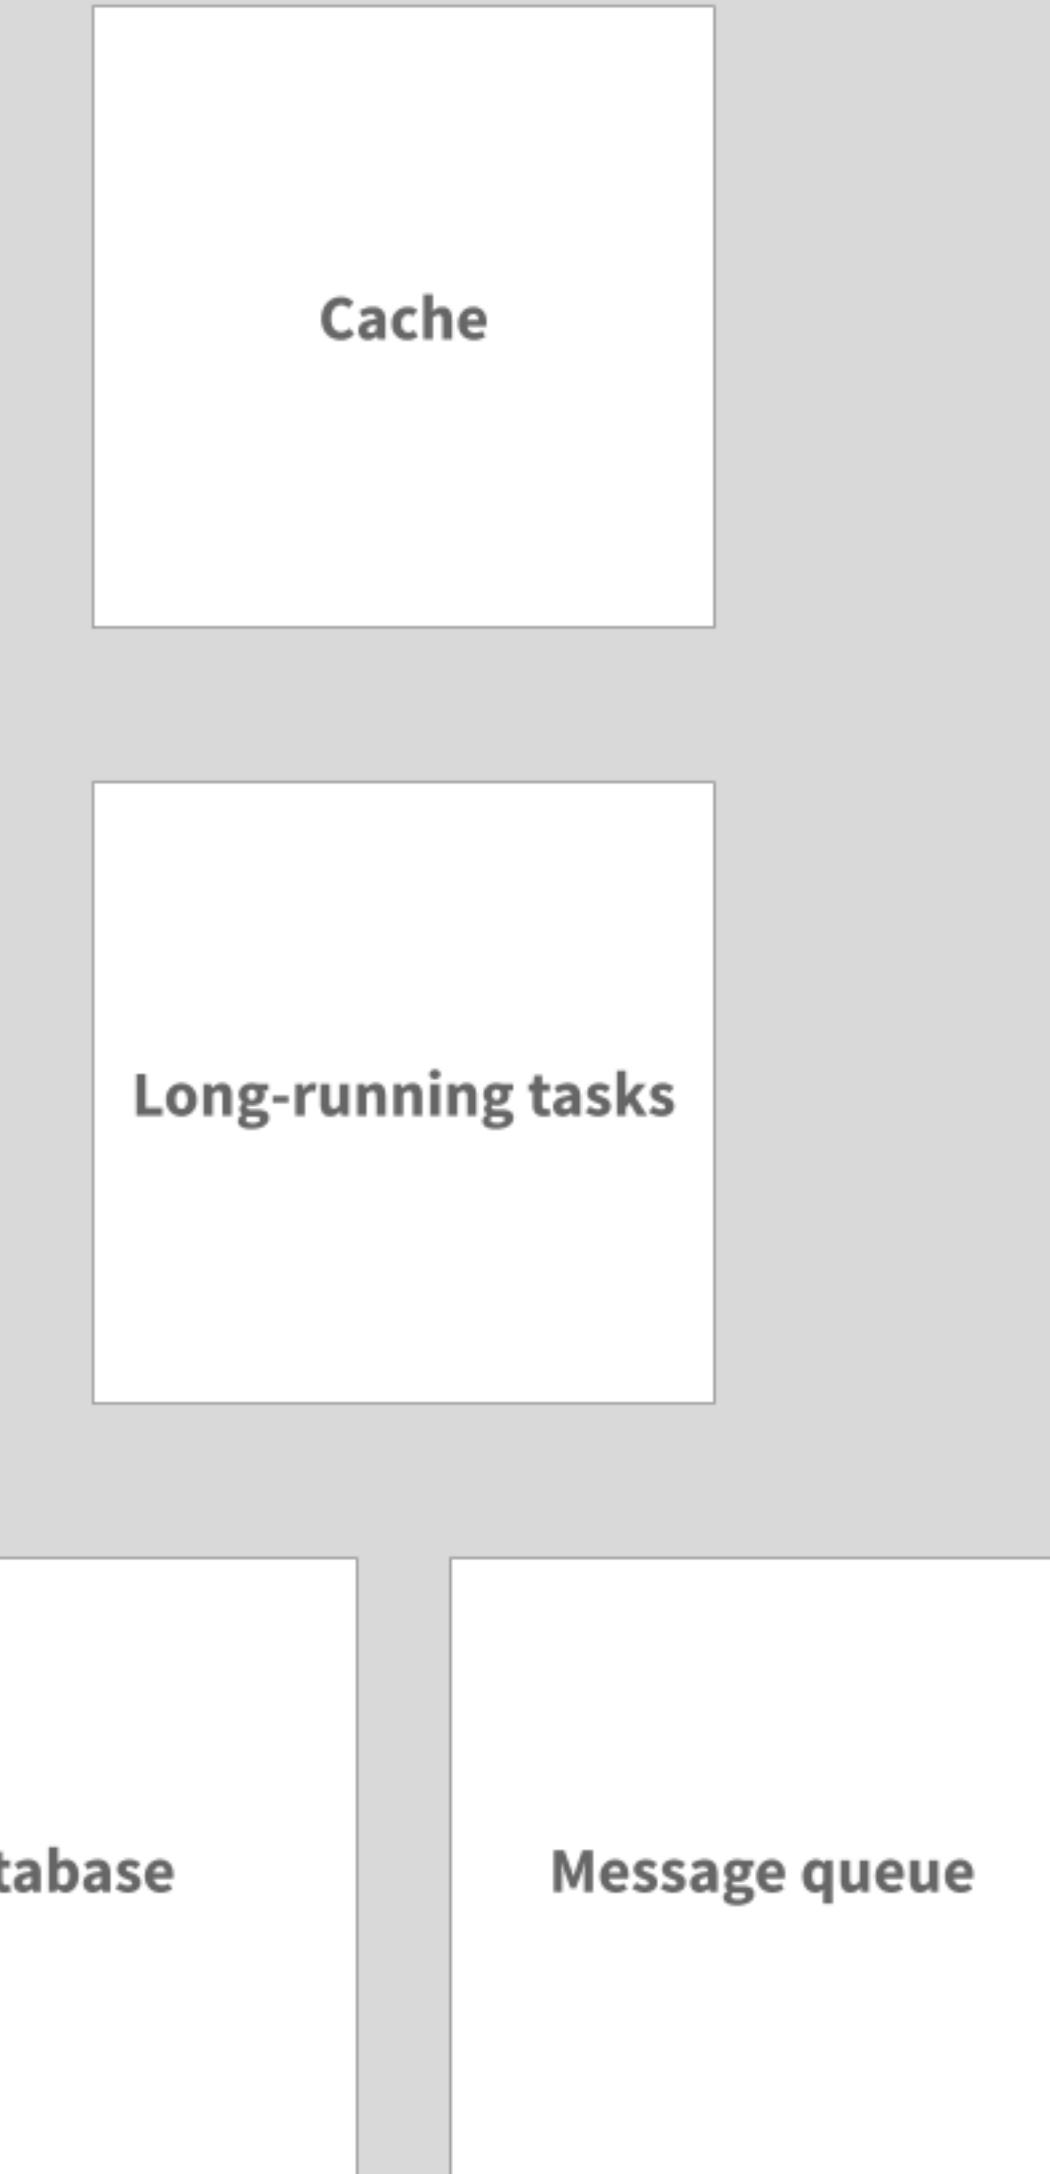
*Out of scope*



# The WSGI stack

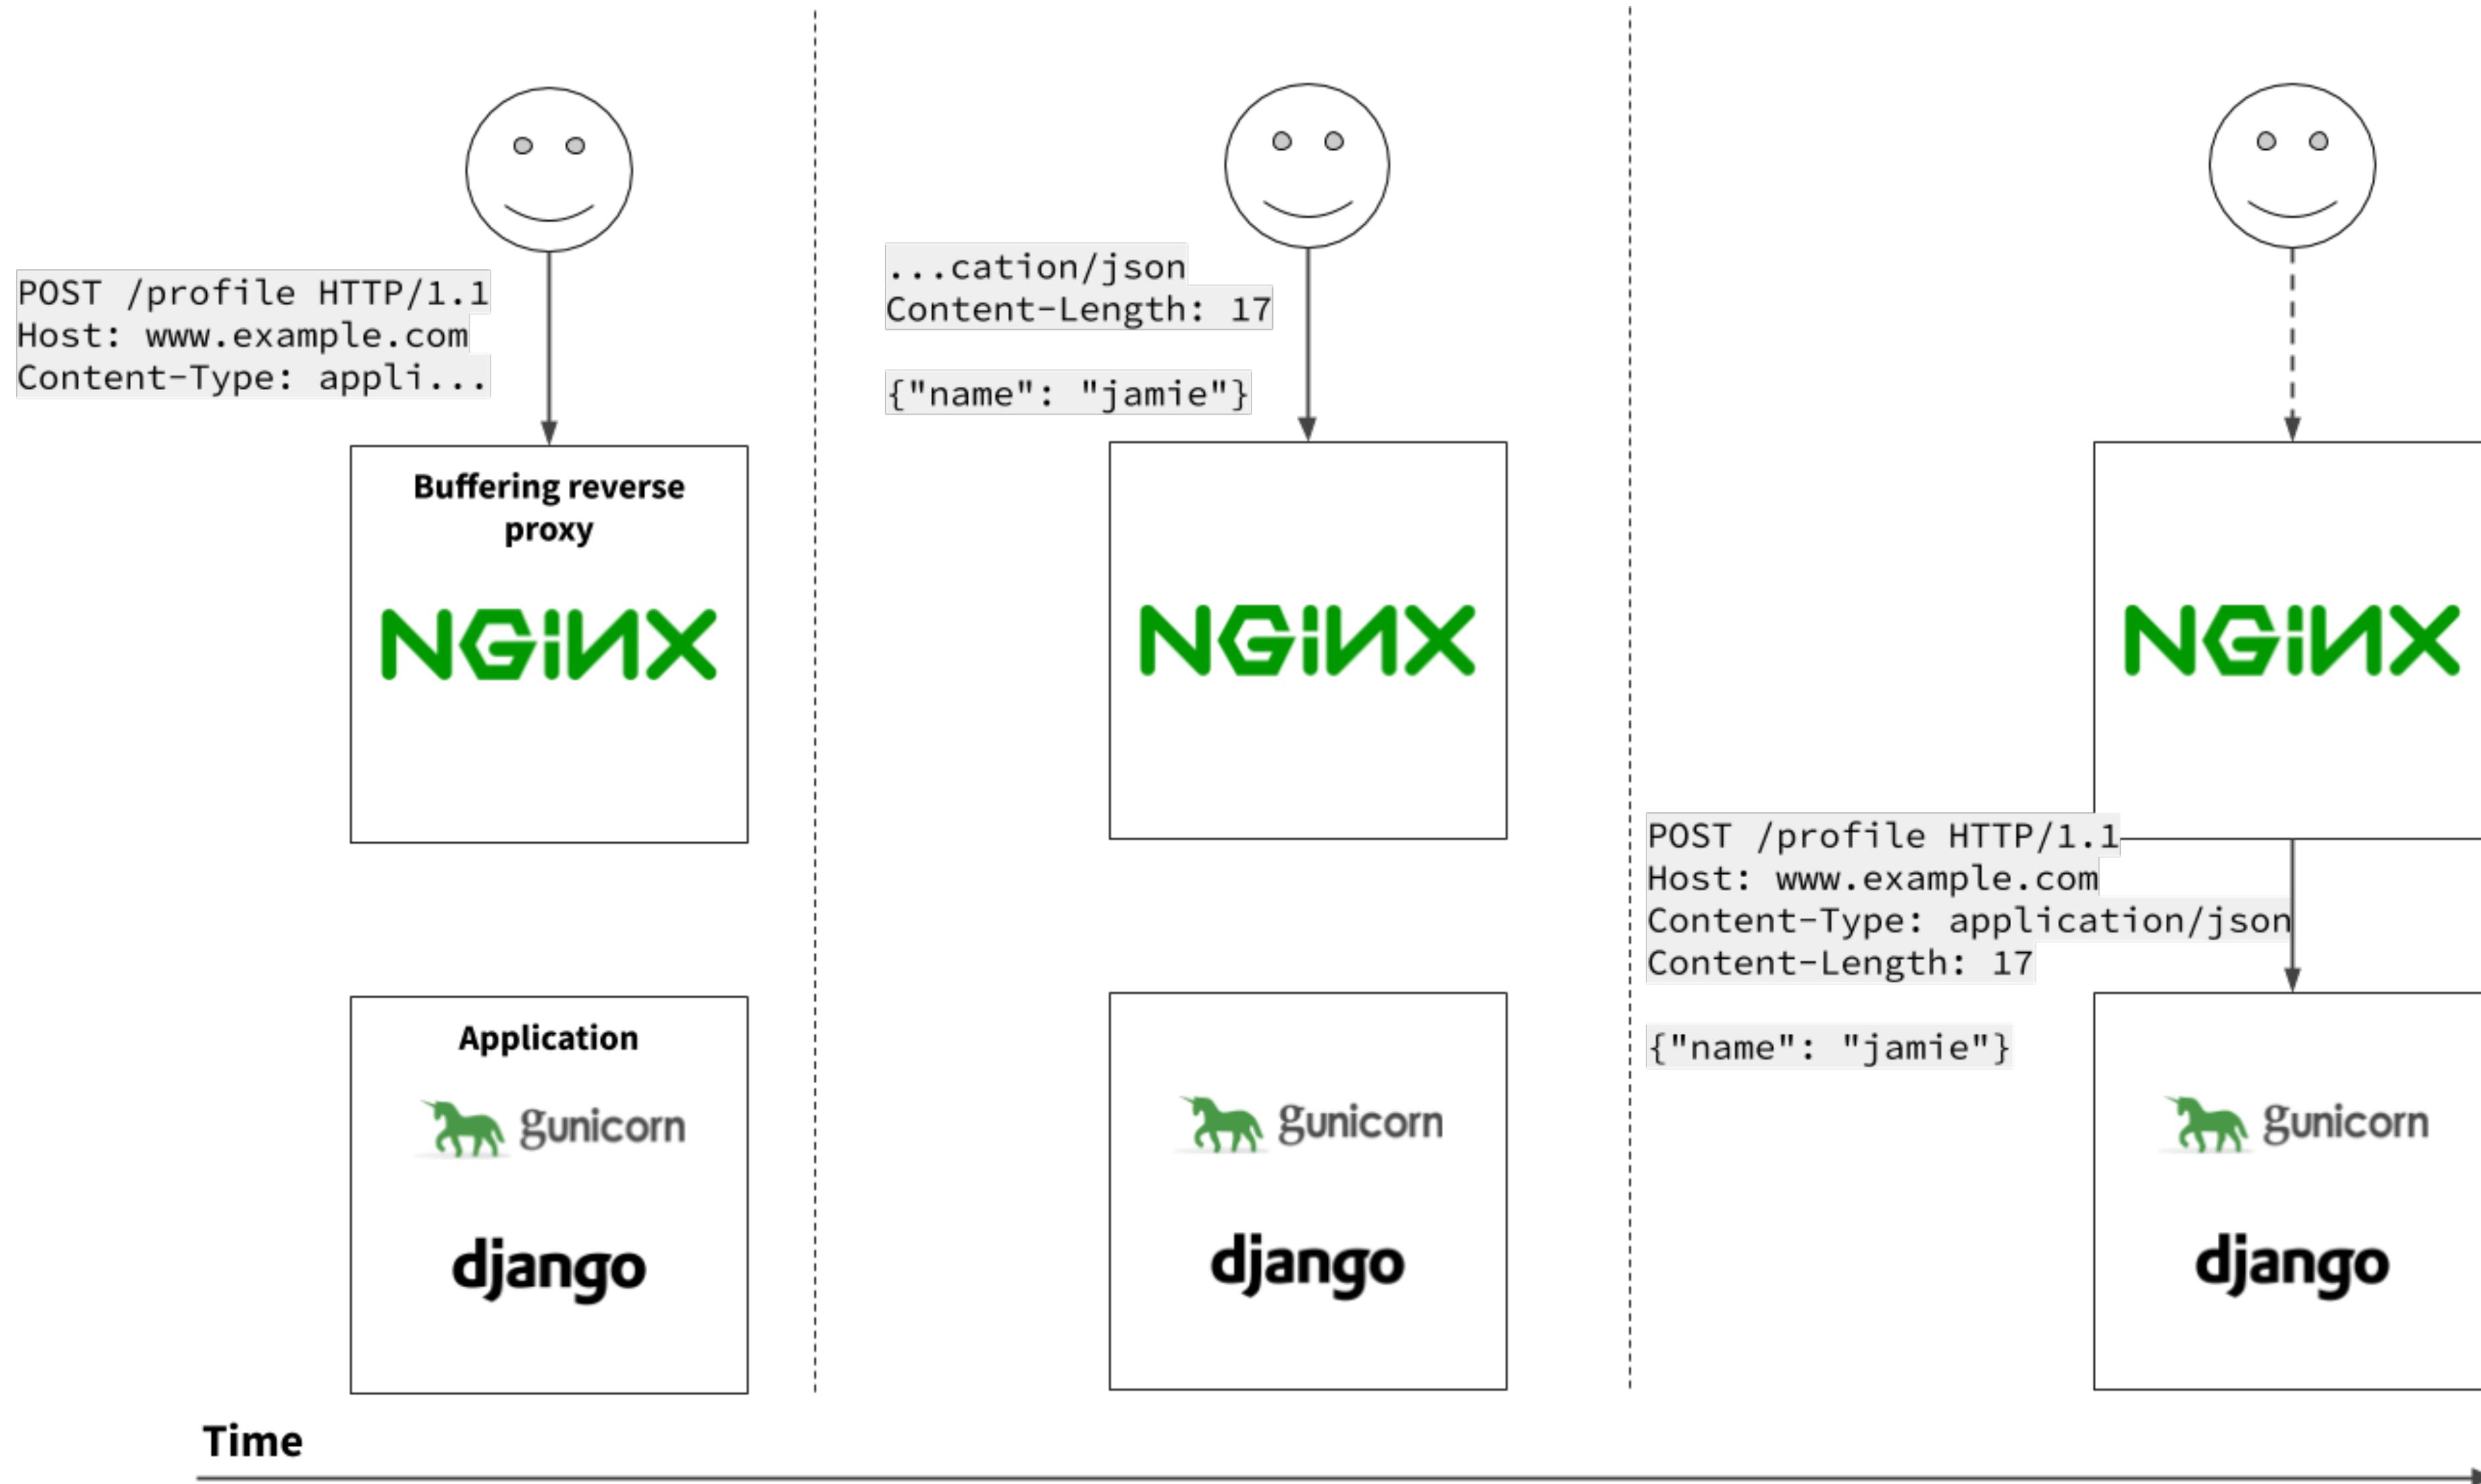


*Out of scope*

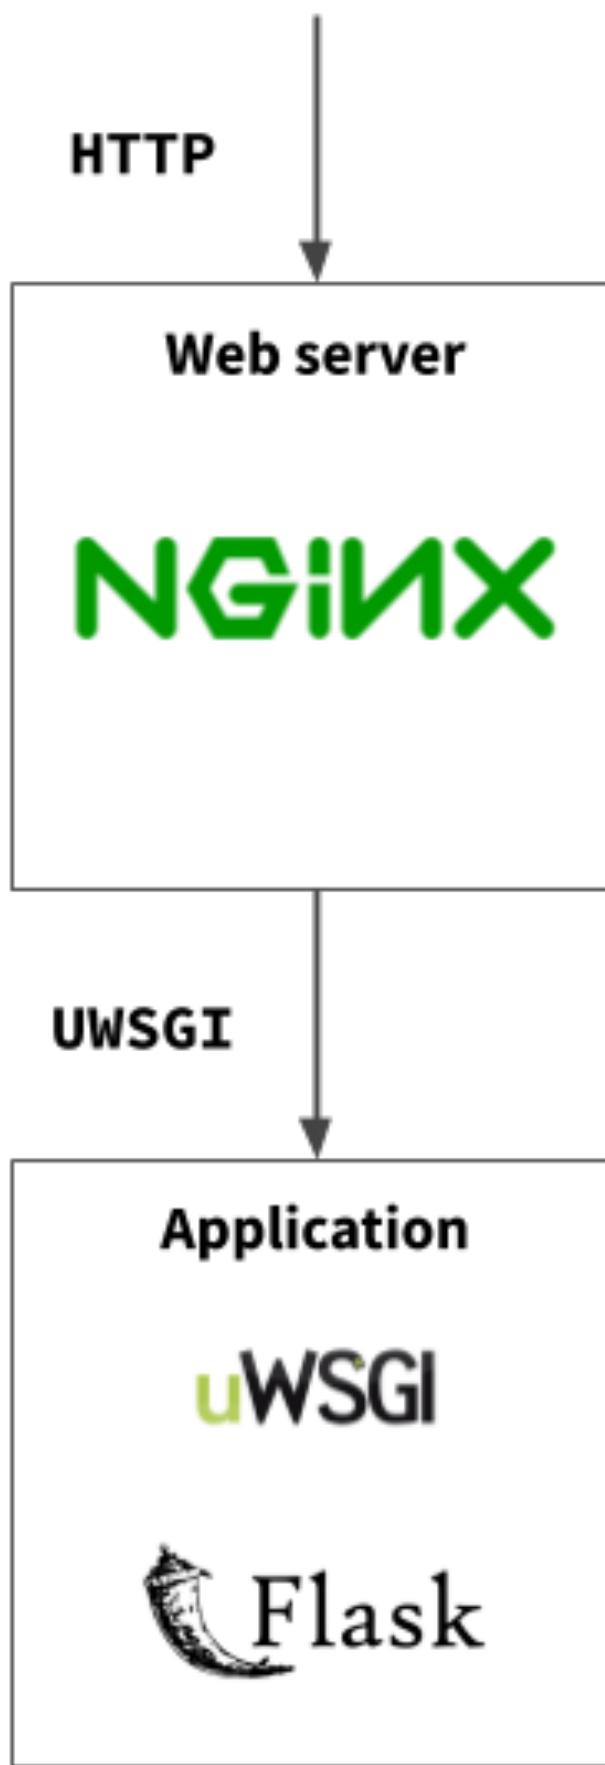




# Gunicorn's design



## Most WSGI servers are designed like this:



the web client

-> ~~the internet~~

-> **the web server** (Nginx)

-> local socket

-> **WSGI server**

-> your application

**The web server...**

- Buffers slow client requests so that app code can stay synchronous (Gunicorn)
- Converts from HTTP to UWSGI protocol (uWSGI)
- “Shields” the Python program

With Kubernetes we can't *expect* there to be a web server

# Buffering reverse proxy

## SOLUTION

### Sidecar proxy in Pod

Run a buffering proxy such as Nginx in the same **Pod** as the application

- ✓ Don't modify the app
- ✓ Maintain similarity with pre-K8S environment
- ✗ Additional container for each Pod
- ✗ More K8S YAML

## ALTERNATIVE

### Remove need for proxy

Run **async WSGI workers** that don't require a buffering proxy

- ✓ No proxy to set up
- ✓ Potentially lower overhead
- ✓ OK with an **Ingress**
- ✗ Potential code changes, app behaviour changes

## GOLD STAR ★

### Service mesh

Use a service mesh to “abstract the network layer”.

- ✓ No proxy specific to app to set up
- ✓ All the pros of a service mesh
- ✗ Have to run a service mesh

# Buffering reverse proxy

## ***RECOMMENDED***

### SOLUTION

#### **Sidecar proxy in Pod**

Run a buffering proxy such as Nginx in the same **Pod** as the application

- ✓ Don't modify the app
- ✓ Maintain similarity with pre-K8S environment
- ✗ Additional container for each Pod
- ✗ More K8S YAML

### ALTERNATIVE

#### **Remove need for proxy**

Run **async WSGI workers** that don't require a buffering proxy

- ✓ No proxy to set up
- ✓ Potentially lower overhead
- ✓ OK with an **Ingress**
- ✗ Potential code changes, app behaviour changes

### GOLD STAR ★

#### **Service mesh**

Use a service mesh to “abstract the network layer”.

- ✓ No proxy specific to app to set up
- ✓ All the pros of a service mesh
- ✗ Have to run a service mesh

# Nginx sidecar

```

containers:
- name: django
  image: jamiehewland/django-app:0aacd1d
  command: ["gunicorn"]
  args: ["--bind=127.0.0.1:8000", "django_app:wsgi"]
  env:
    - name: DJANGO_SETTINGS_MODULE
      value: django_app.settings.production

- name: nginx
  image: nginx:stable
  ports:
    - containerPort: 80
  volumeMounts:
    - name: nginx-config
      mountPath: /etc/nginx/conf.d

volumes:
- name: nginx-config
  configMap:
    name: nginx-config

```



```

apiVersion: v1
kind: ConfigMap
metadata:
  name: nginx-config
data:
  wsgi.conf: |
    upstream wsgi {
      server 127.0.0.1:8000;
    }

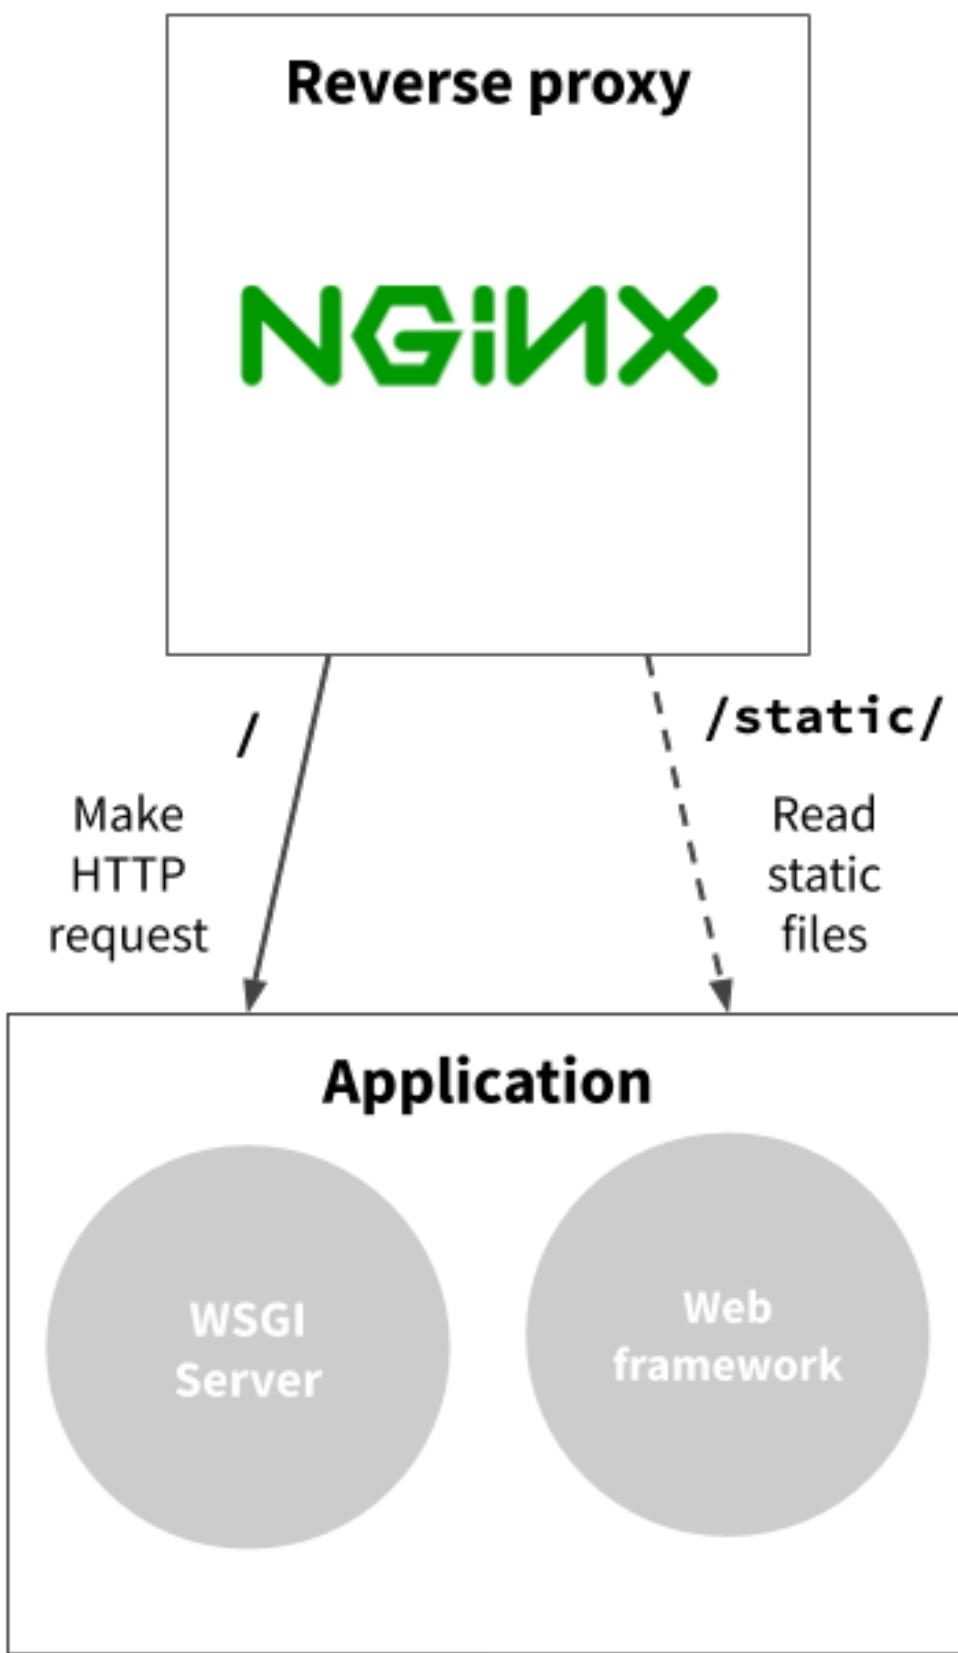
    server {
      listen 80;
      server_name _;

      location / {
        proxy_set_header X-Forwarded-For $proxy_add_x_forwarded_for;
        proxy_set_header X-Forwarded-Proto $scheme;
        proxy_set_header Host $http_host;
        proxy_redirect off;
        proxy_pass http://wsgi;
      }
    }

```



# Serving static files



```
location /static/ {  
    alias /path/to/static/files;  
}
```

```
location / {  
    proxy_set_header ...  
    proxy_pass http://127.0.0.1:8000;  
}
```

GET /admin/

GET /static/admin/css/base.css

[Django]  
[Nginx]

## uWSGI:

*[...] it's inefficient to serve static files via uWSGI. Instead, serve them directly from Nginx and completely bypass uWSGI."*

# Serving static files

## SOLUTION

### Sidecar web server in Pod

Run a web server such as Nginx in the same **Pod** as the application, share a **Volume** with static files

- ✓ Don't modify the app
- ✓ Maintain similarity with pre-K8S environment
- ✗ Possibly additional container for each **Pod**
- ✗ **Volume** increases startup time

## ALTERNATIVE

### Serve static files from S3

Use **django-storages** or similar library to upload static files to S3 or similar

- ✓ Simple & easy
- ✗ Additional step to upload files
- ✗ Unoptimised (compression/caching headers)

## GOLD STAR ★

### WhiteNoise + CDN

Serve optimised static files with Python using WhiteNoise library

- ✓ Automatically compressed assets
- ✓ Smart caching headers
- ✓ App & static files always in sync
- ✗ Potentially larger container images

# Serving static files

## SOLUTION

### Sidecar web server in Pod

Run a web server such as Nginx in the same **Pod** as the application, share a **Volume** with static files

- ✓ Don't modify the app
- ✓ Maintain similarity with pre-K8S environment
- ✗ Possibly additional container for each **Pod**
- ✗ **Volume** increases startup time

## ALTERNATIVE

### Serve static files from S3

Use **django-storages** or similar library to upload static files to S3 or similar

- ✓ Simple & easy
- ✗ Additional step to upload files
- ✗ Unoptimised (compression/caching headers)

## ***RECOMMENDED***

### GOLD STAR ★

### WhiteNoise + CDN

Serve optimised static files with Python using WhiteNoise library

- ✓ Automatically compressed assets
- ✓ Smart caching headers
- ✓ App & static files always in sync
- ✗ Potentially larger container images

# WhiteNoise

## Advantages

- Serve static files directly from Python
- Compressed assets (w/correct headers): GZIP, Brotli
- Automatic far-future cache headers

## Isn't serving static files from Python horribly slow?

- Use a CDN
- All assets generated offline
- WSGI servers can use the sendfile syscall

<http://whitenoise.evans.io/en/stable/>

› pip install whitenoise

## Django:

1. Add middleware class
2. (Optional) Add staticfiles storage class
3. Remember to run collectstatic

## Flask:

1. Add WSGI middleware
2. (Optional) Configure static file directories



# Deployments

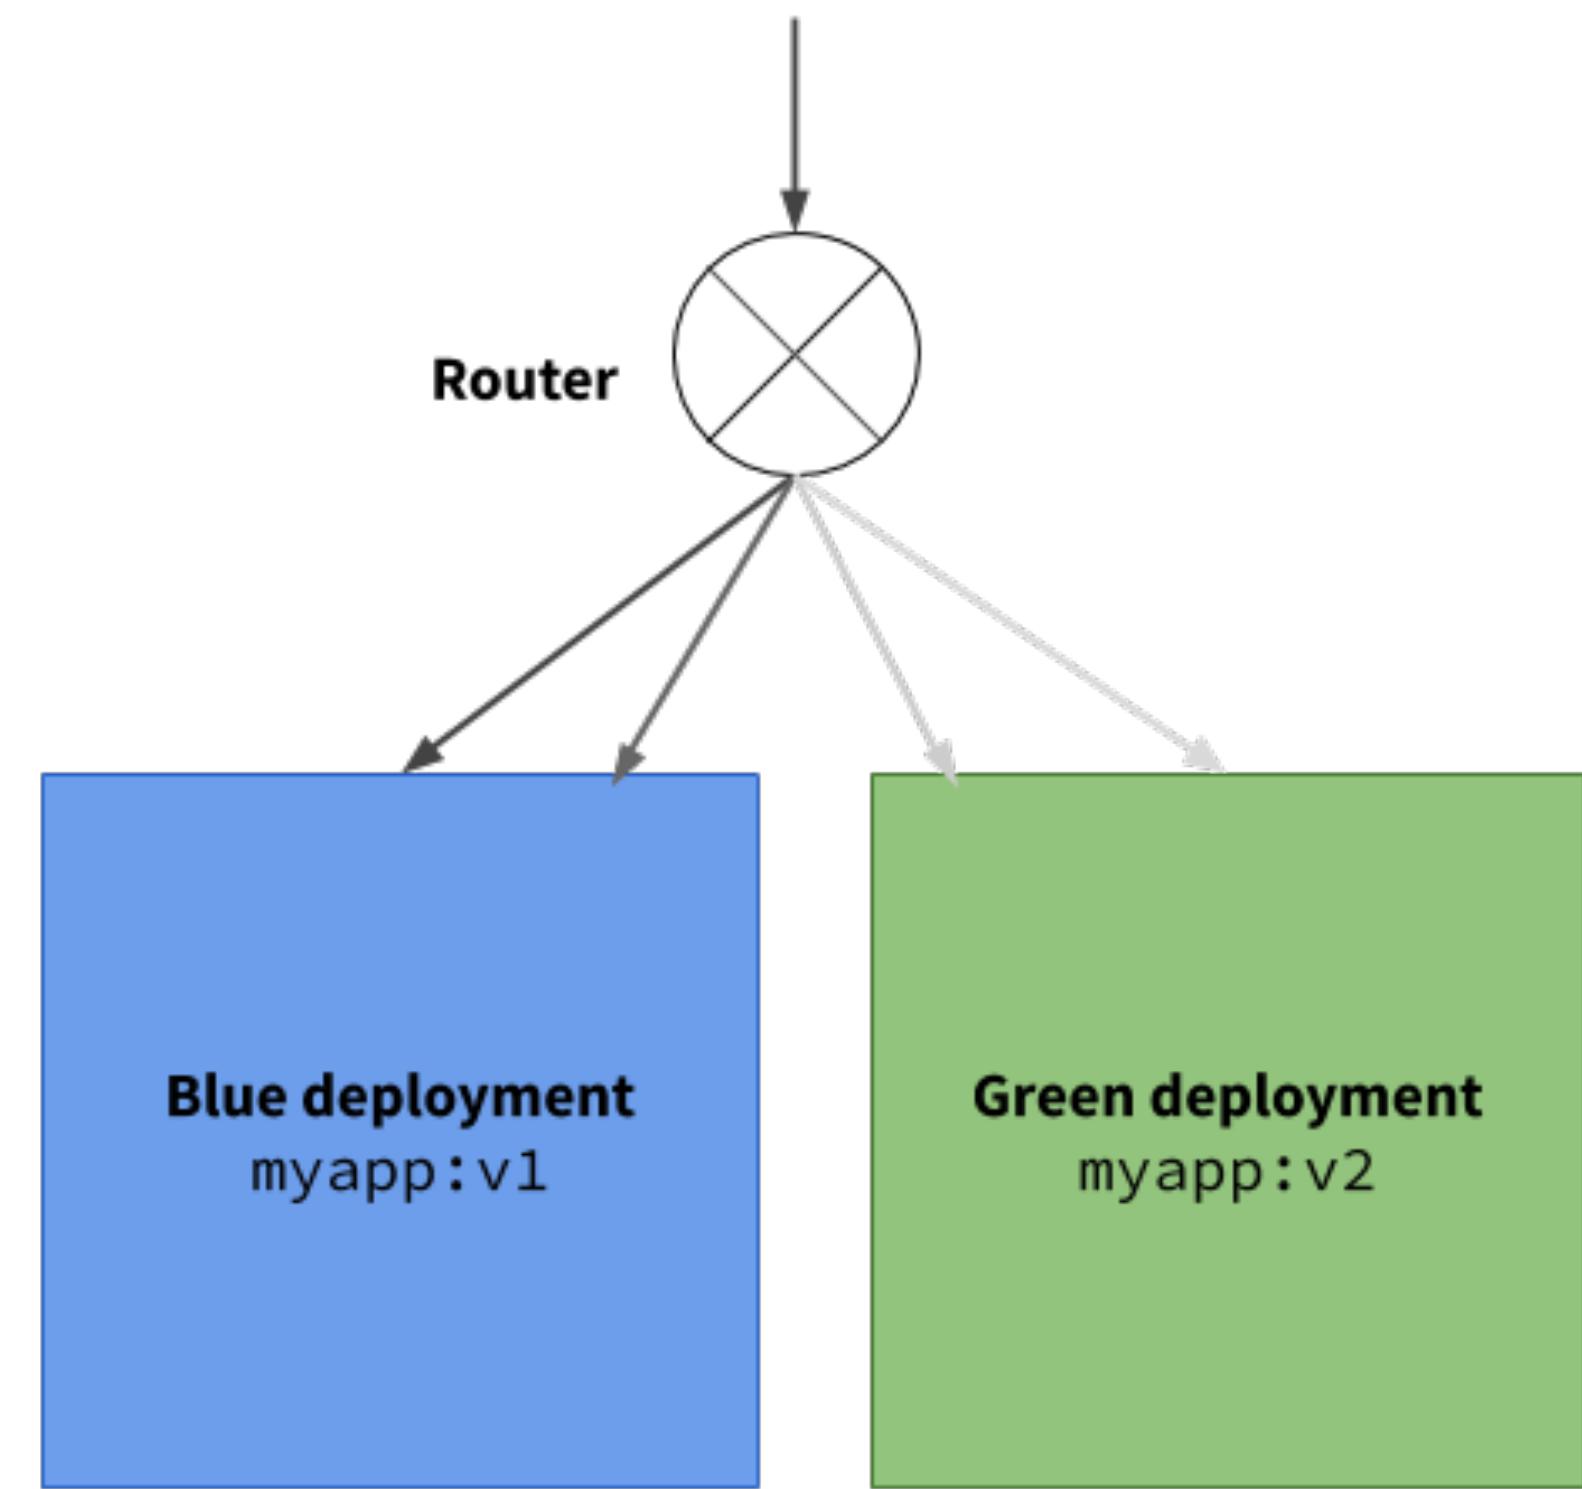
***Previously:*** update the code, then:

- Gunicorn and uWSGI support graceful reload via SIGHUP
- Some may just restart the process

**With containers we *never* change the running code,  
rather:**

- Build a new image with the new code
- Run new container with new image
- Switch traffic over to new container
- Shut down old container

<https://martinfowler.com/bliki/BlueGreenDeployment.html>



# Deployments

## SOLUTION

### Deployment

A **Deployment** resource with default settings such as RollingUpdate.

- ✓ Basic blue/green deploy
- ✗ Limited deployment history
- ✗ Limited update strategies

## NEXT STEPS

### Multiple Deployments

Shift traffic between two or more Deployments

- ✓ Fast roll-backs
- ✓ Open the door to more advanced deployments
- ✗ Increased resource usage (run 2 copies)
- ✗ Need extra tool to manage deployment

## GOLD STAR ★

### Advanced deployment strategies

Canarying, gradual red/black deployments, traffic shadowing

- ✓ “Test in production!”
- ✗ Requires sophisticated tooling & expertise

# Deployments

## ***RECOMMENDED***

### SOLUTION

## Deployment

A **Deployment** resource with default settings such as RollingUpdate.

- ✓ Basic blue/green deploy
- ✗ Limited deployment history
- ✗ Limited update strategies

### NEXT STEPS

## Multiple Deployments

Shift traffic between two or more Deployments

- ✓ Fast roll-backs
- ✓ Open the door to more advanced deployments
- ✗ Increased resource usage (run 2 copies)
- ✗ Need extra tool to manage deployment

### GOLD STAR ★

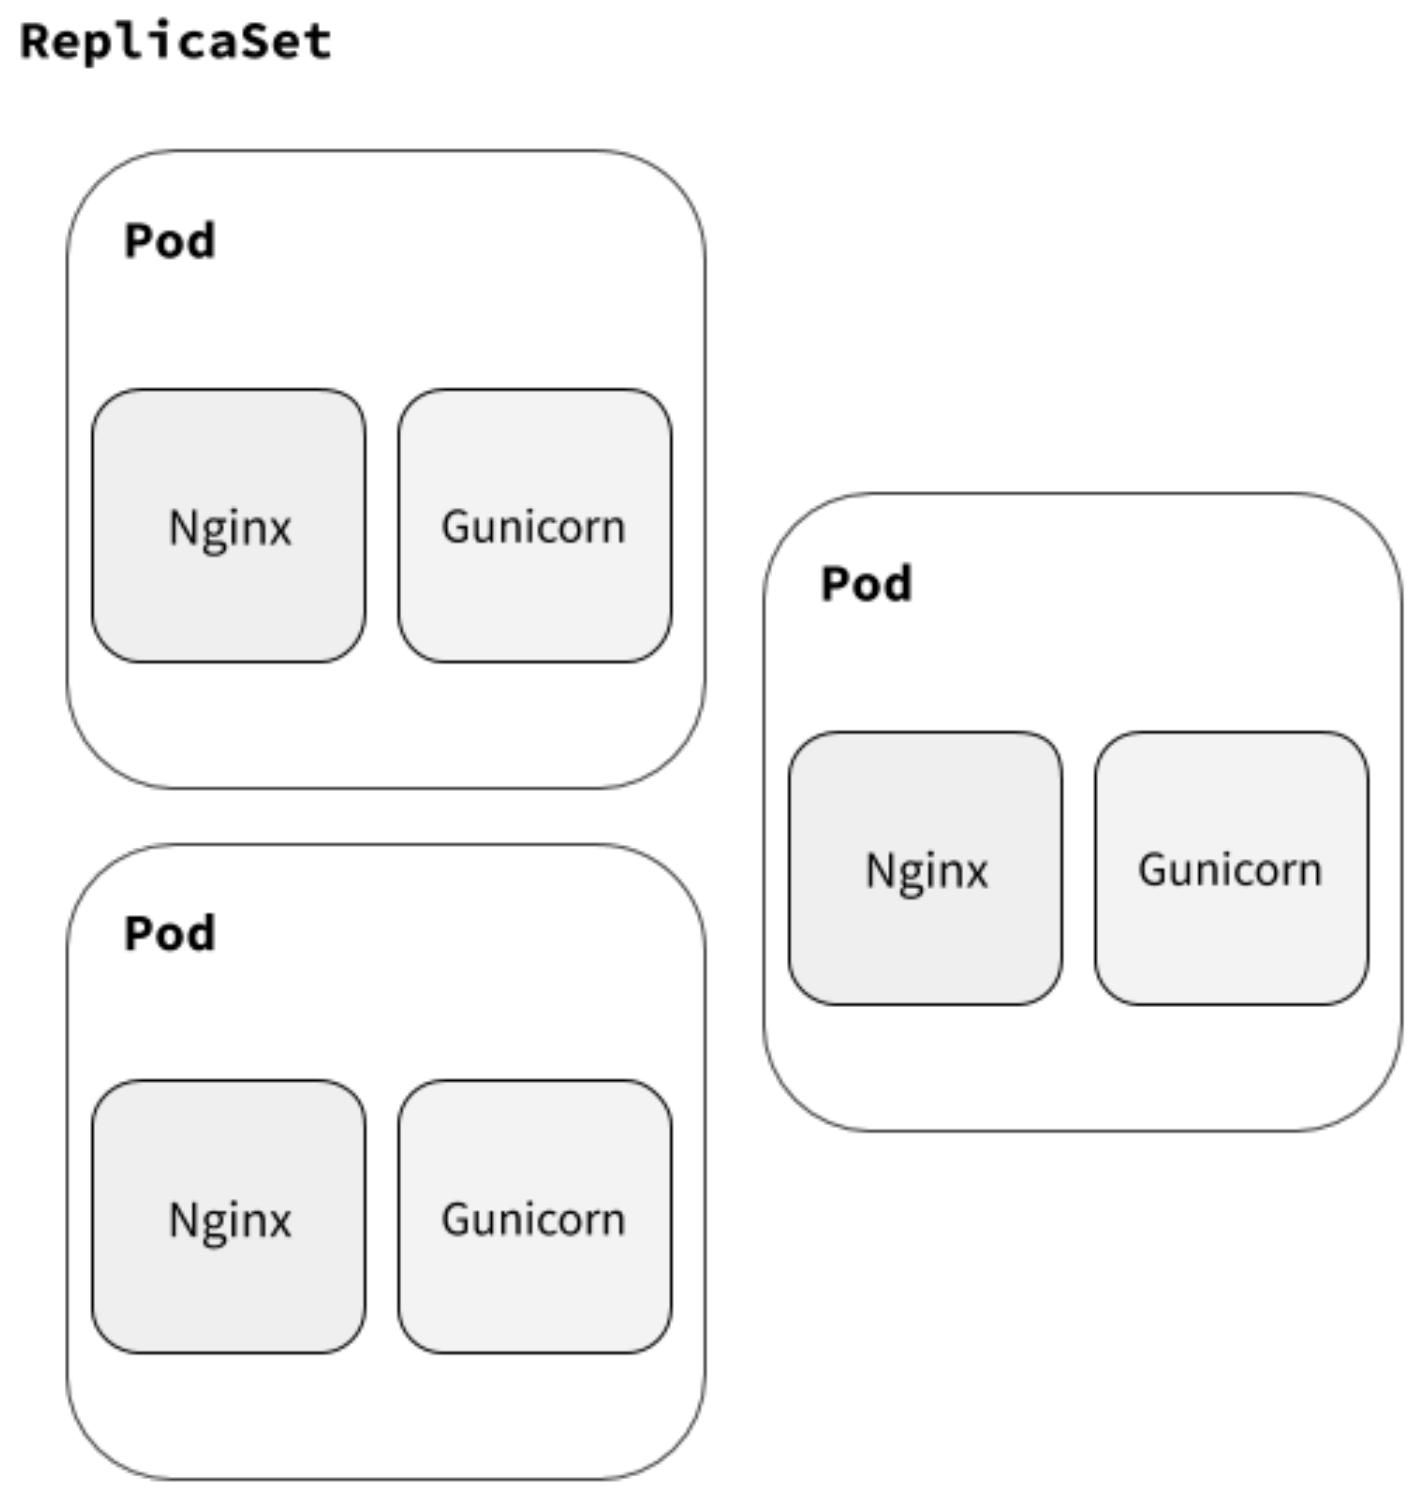
## Advanced deployment strategies

Canarying, gradual red/black deployments, traffic shadowing

- ✓ “Test in production!”
- ✗ Requires sophisticated tooling & expertise

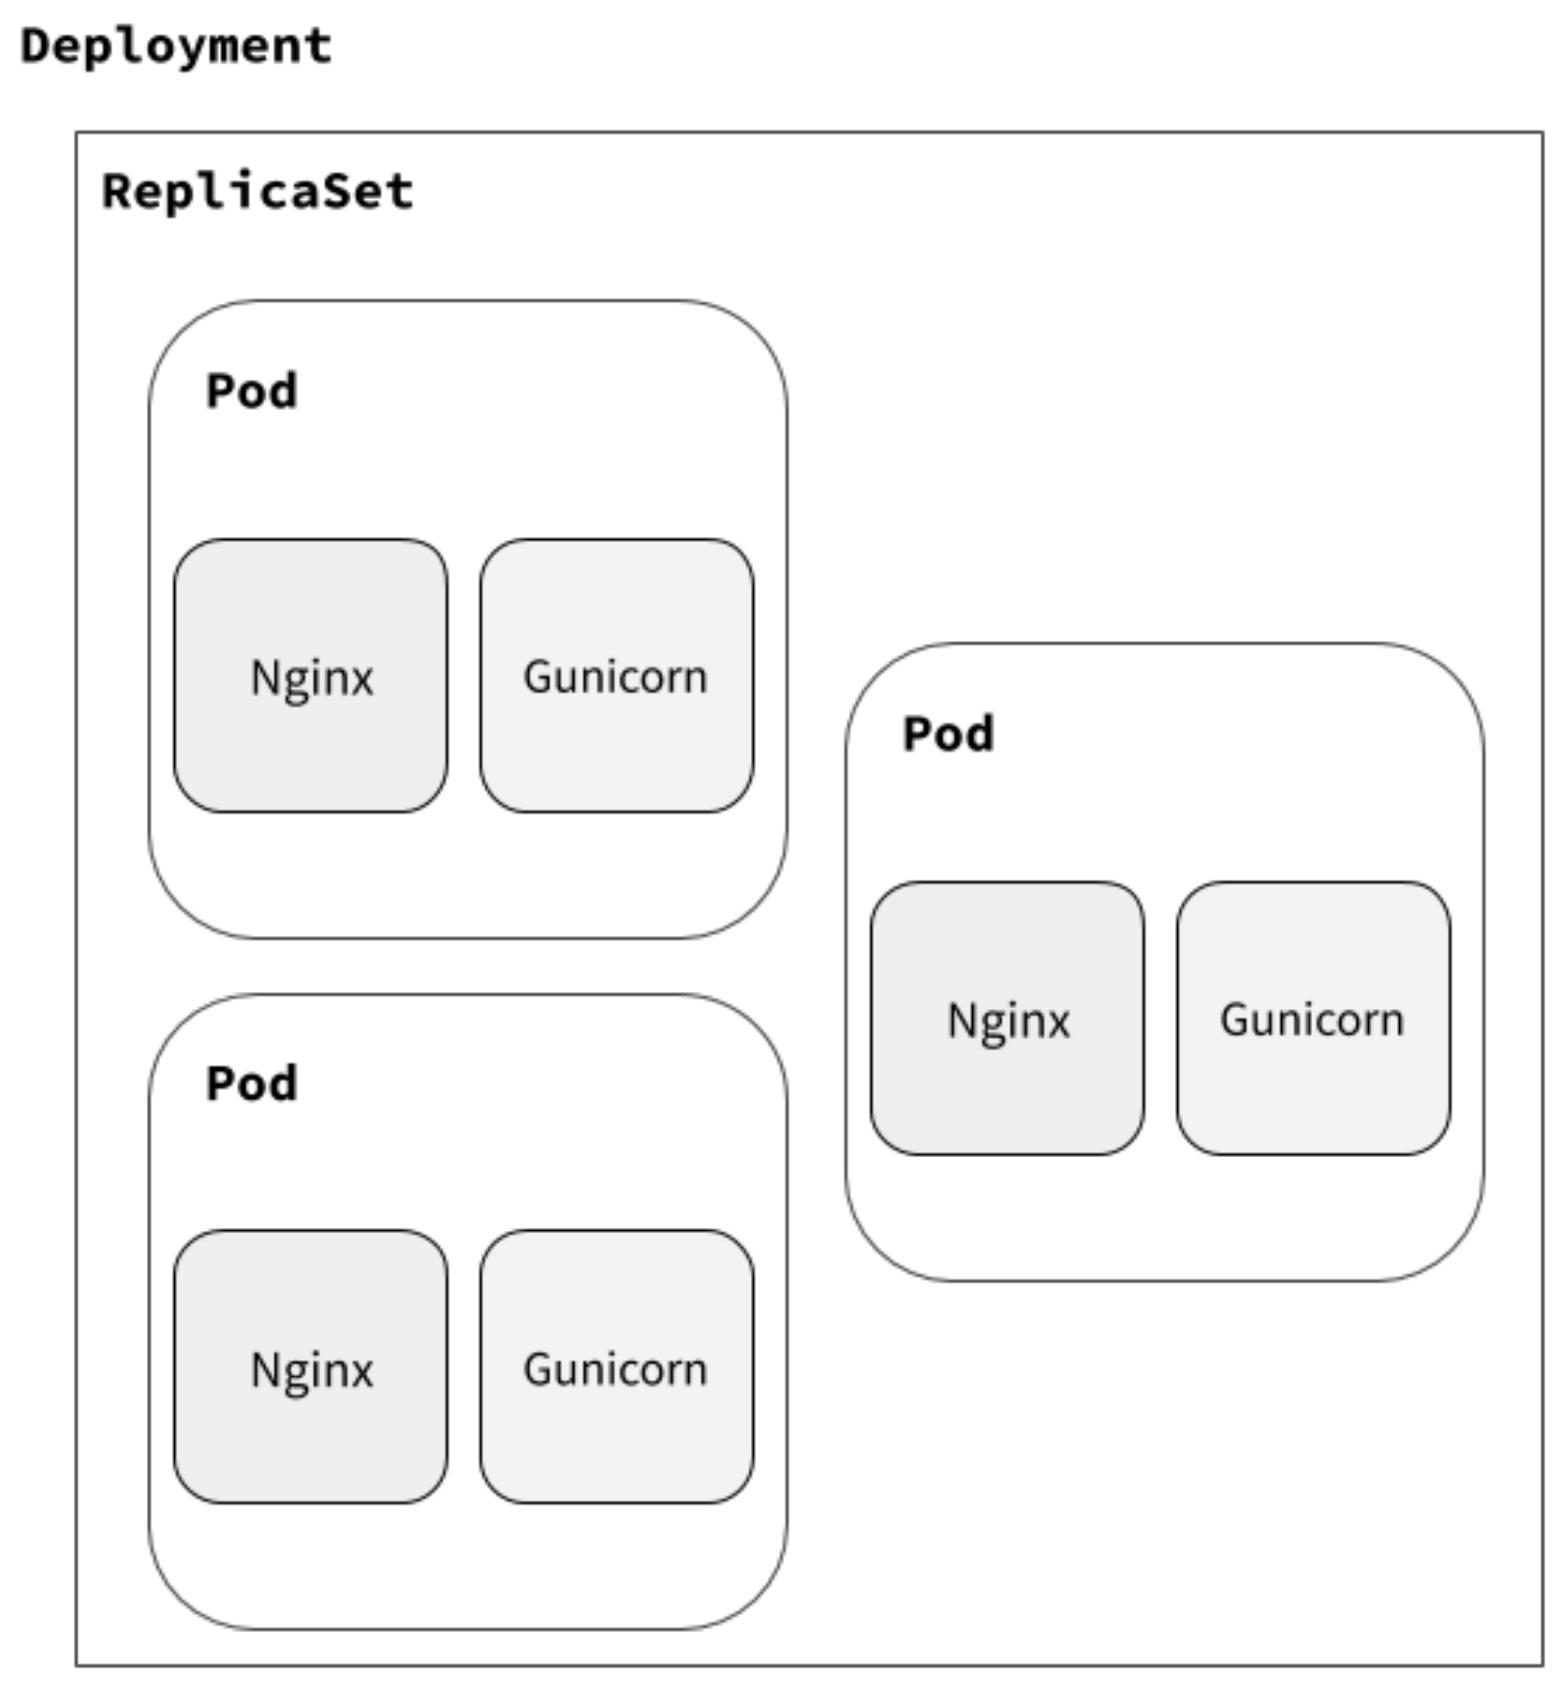
# Kubernetes Deployment

## Deployment



```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: django-deployment
spec:
  replicas: 3
  selector:
    matchLabels:
      app: django-app
  template:
    metadata:
      labels:
        app: django-app
    spec:
      containers:
        - name: django
          image: jamiehewland/django-app:0aacd1d
          env:
            - name: DJANGO_SETTINGS_MODULE
              value: django_app.settings.production
        - name: nginx
          image: nginx:stable
          ports:
            - containerPort: 80
          volumeMounts:
            - name: nginx-config
              mountPath: /etc/nginx/conf.d
          volumes:
            - name: nginx-config
          configMap:
            name: nginx-config
```

# Kubernetes Deployment



3 replicas

Django container

Nginx container

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: django-deployment
spec:
  replicas: 3
  selector:
    matchLabels:
      app: django-app
  template:
    metadata:
      labels:
        app: django-app
    spec:
      containers:
        - name: django
          image: jamiehewland/django-app:0aacd1d
          env:
            - name: DJANGO_SETTINGS_MODULE
              value: django_app.settings.production
        - name: nginx
          image: nginx:stable
          ports:
            - containerPort: 80
          volumeMounts:
            - name: nginx-config
              mountPath: /etc/nginx/conf.d
          volumes:
            - name: nginx-config
          configMap:
            name: nginx-config
```

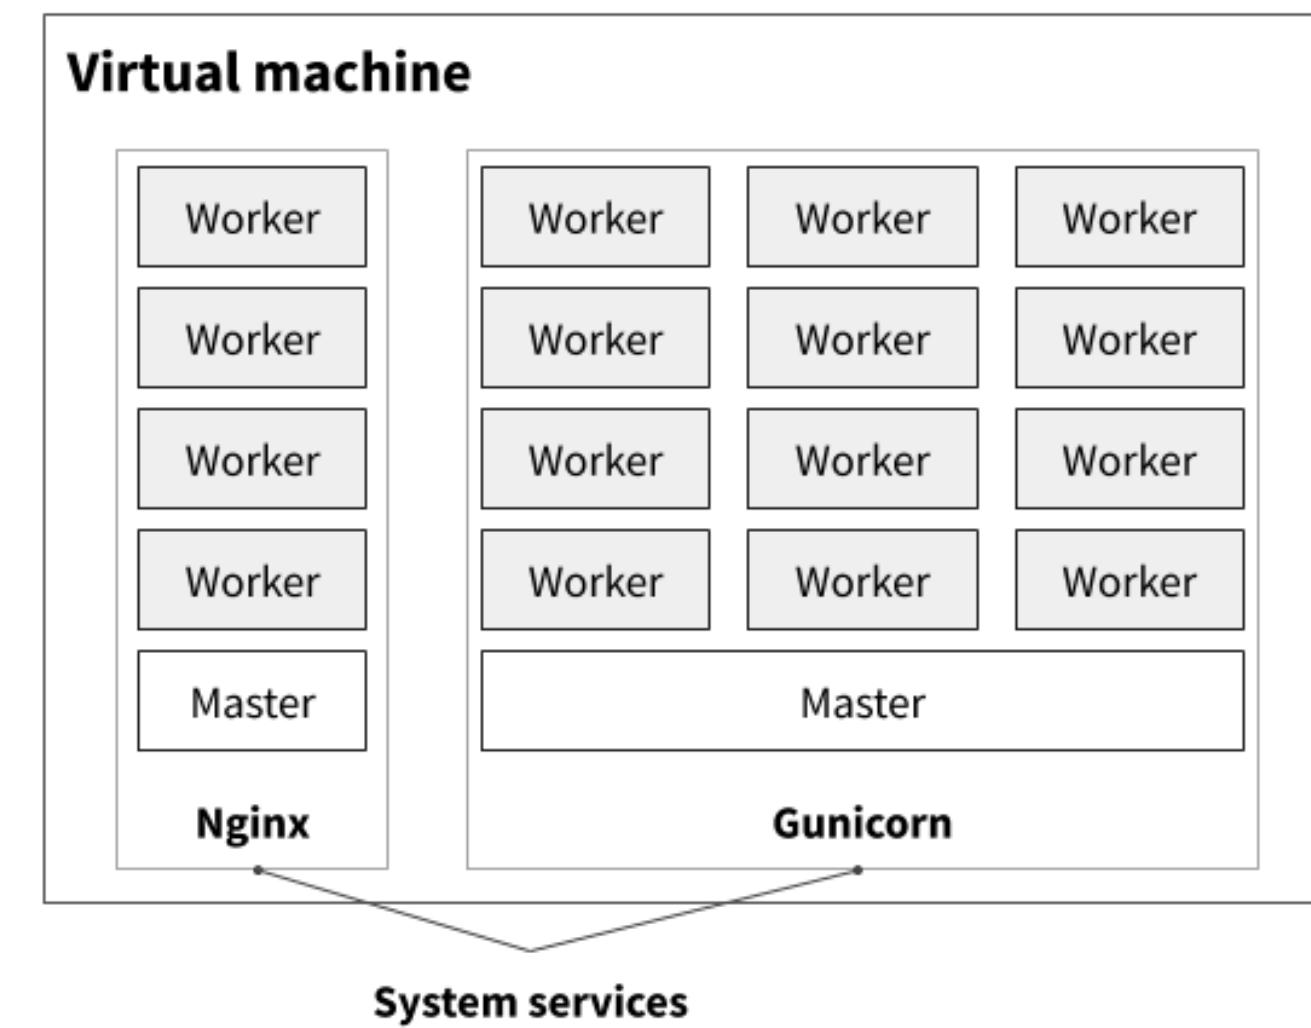
# Smaller instances

## Pros:

- Finer-grained scaling
- More redundancy
- Better for canarying

## Cons:

- Increased overhead (especially RAM)
- Greater need for automation



Virtual machines  $\blacktriangleleft \triangleright$  Kubernetes



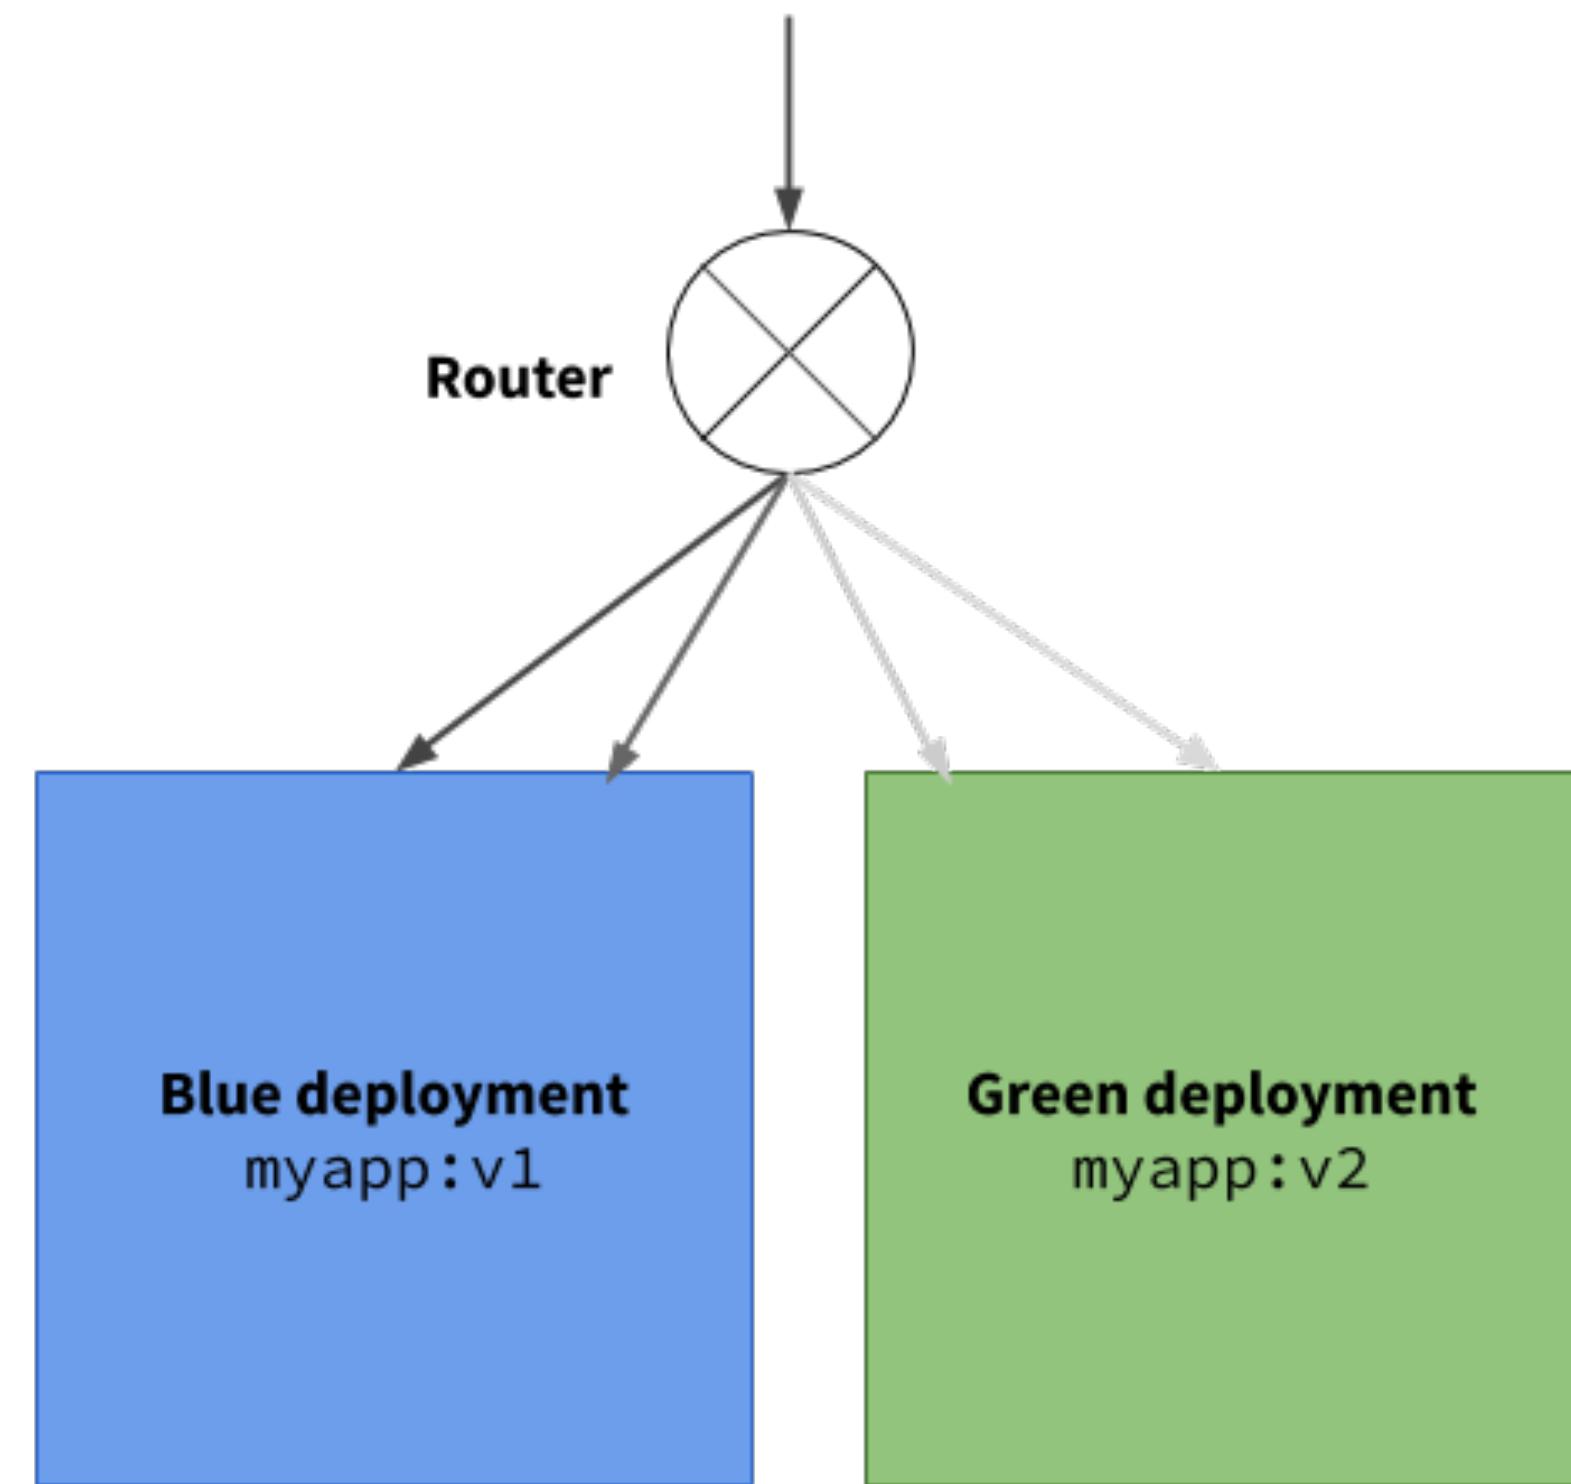
# Configuration

***Previously:*** update the config by...

- Running Configuration Management tool (e.g. Chef, Puppet, Ansible)
- Reload/restart server

**Like code, we *never* want the config a container runs with to change:**

- Run new container with the new config
- Switch traffic over to new container
- Shut down old container



# Configuration

## STARTING OUT

### Environment variables

Set environment variables for containers in the **Deployment**

- ✓ It works
- ✗ Lots of YAML
- ✗ Have to pack everything into strings

## BETTER

### ConfigMap

Store config in a ConfigMap and import in app via env vars or file volume

- ✓ Separate app & config
- ✗ ConfigMaps are mutable

## BEST PRACTICE

### Immutable ConfigMaps

Create a new ConfigMap for each config change

- ✓ Keep record of previous configs
- ✓ Not difficult to do (use Kustomize or Helm)
- ✗ Probably need additional tool

# Configuration

## STARTING OUT

### Environment variables

Set environment variables for containers in the **Deployment**

- ✓ It works
- ✗ Lots of YAML
- ✗ Have to pack everything into strings

## BETTER

### ConfigMap

Store config in a **ConfigMap** and import in app via env vars or file volume

- ✓ Separate app & config
- ✗ **ConfigMaps** are mutable

## RECOMMENDED

## BEST PRACTICE

### Immutable ConfigMaps

Create a new **ConfigMap** for each config change

- ✓ Keep record of previous configs
- ✓ Not difficult to do (use Kustomize or Helm)
- ✗ Probably need additional tool

# ConfigMap



```
apiVersion: v1
kind: ConfigMap
metadata:
  name: django-config
data:
  LOG_LEVEL: info
  DATABASE_HOST: django-app.cvqdtcqjj6f2.us-west-2.rds.amazonaws.com
  DATABASE_NAME: django-app
```

```
containers:
- name: django
  image: jamiehewland/django-app:0aacd1d
  env:
    - name: DJANGO_SETTINGS_MODULE
      value: django_app.settings.production
  envFrom:
  - configMapRef:
      name: django-config
```



# Credentials

## STARTING OUT

### Store with the configuration

Store credentials just like you would any other configuration

- ✓ Possibly no worse than what you're doing already
- ✗ Credentials stored in plain-text in etcd

## BEST PRACTICE

### Secret

**Store credentials in a Secret and import in app via env vars or (preferably) file volume**

- ✓ Separate sensitive secrets from config
- ✗ May have to create **Secrets** out-of-band
- ✗ Make sure your K8S is configured to use envelope encryption

## GOLD STAR ★

### HashiCorp Vault or similar

Store secrets (possibly dynamic ones) in a dedicated service

- ✓ About as secure as it gets right now
- ✓ Other advantages to Vault, e.g. auditing
- ✗ Additional infrastructure
- ✗ Likely requires app-level changes

# Credentials

## STARTING OUT

### Store with the configuration

Store credentials just like you would any other configuration

- ✓ Possibly no worse than what you're doing already
- ✗ Credentials stored in plain-text in etcd

#### ***RECOMMENDED***

#### BEST PRACTICE

### Secret

Store credentials in a **Secret** and import in app via env vars or (preferably) file volume

- ✓ Separate sensitive secrets from config
- ✗ May have to create **Secrets** out-of-band
- ✗ Make sure your K8S is configured to use envelope encryption

#### GOLD STAR ★

### HashiCorp Vault or similar

Store secrets (possibly dynamic ones) in a dedicated service

- ✓ About as secure as it gets right now
- ✓ Other advantages to Vault, e.g. auditing
- ✗ Additional infrastructure
- ✗ Likely requires app-level changes

# Secret



```
> kubectl create secret generic django-credentials \
    --from-literal DATABASE_USER="$DATABASE_USER" \
    --from-literal DATABASE_PASSWORD="$DATABASE_PASSWORD"
```

```
containers:
- name: django
  image: jamiehewland/django-app:0aacd1d
  env:
  - name: DJANGO_SETTINGS_MODULE
    value: django_app.settings.production
  envFrom:
  - configMapRef:
      name: django-config
  - secretRef:
      name: django-credentials
```



# Database migrations

## STARTING OUT

### **initContainers**

Add a container in the **initContainers** section to run on **Pod** startup

- Basic automation
- X Run on every **Pod** start
- X No post-/pre-deployment control

## BEST PRACTICE

### **Job triggered manually**

**Run a one-off Job manually each time a migration needs to be run**

- Single process running migration
- Separate deployment of schema changes from deployment of app
- X Automation?

## GOLD STAR

### **Job triggered by CD**

Post- and/or pre-deployment **Job** run automatically by Continuous Deployment system

- Automated
- X Requires mature CD system

# Database migrations

## STARTING OUT

### **initContainers**

Add a container in the `initContainers` section to run on **Pod** startup

- Basic automation
- X Run on every **Pod** start
- X No post-/pre-deployment control

### ***RECOMMENDED***

#### BEST PRACTICE

##### **Job triggered manually**

**Run a one-off Job manually each time a migration needs to be run**

- Single process running migration
- Separate deployment of schema changes from deployment of app
- X Automation?

#### GOLD STAR ★

##### **Job triggered by CD**

Post- and/or pre-deployment **Job** run automatically by Continuous Deployment system

- Automated
- X Requires mature CD system

## Job migrations

Run a migration using the image for the app

```
apiVersion: batch/v1
kind: Job
metadata:
  name: django-migrations
spec:
  template:
    spec:
      containers:
        - name: django
          image: jamiehewland/django-app:0aacd1d
          command: ["python", "manage.py", "migrate"]
          env:
            - name: DJANGO_SETTINGS_MODULE
              value: django_app.settings.production
            - name: DATABASE_USER
              valueFrom:
                secretKeyRef:
                  name: django-credentials
                  key: DATABASE_USER
            - name: DATABASE_PASSWORD
              valueFrom:
                secretKeyRef:
                  name: django-credentials
                  key: DATABASE_PASSWORD
  restartPolicy: Never
```

<b>Problem</b>	<b>“Legacy” solution</b>	<b>Recommended Kubernetes solution</b>
Buffering reverse proxy	Nginx on host	Pod with sidecar proxy
Static file serving	Nginx on host	WhiteNoise + CDN
Deployments	WSGI server reload signal	<a href="#">Deployment</a> with blue/green strategy
Configuration	Configuration Management tool	<a href="#">ConfigMap</a> (and/or 12-factor app)
Credentials	Configuration Management tool	<a href="#">Secret</a>
Database migrations	Manual or CD step	<a href="#">Job migration container</a>

tl;dl

This is a lot to get right (or wrong)

Why is it so complicated?

- Established tools like Gunicorn and uWSGI were designed before containers/Kubernetes
  - Process management
  - Hot-code reload
  - User-switching
  - Daemonizing
  - Many tuneable settings

```
uwsgi --chdir=/path/to/your/project \
--module=mysite.wsgi:application \
--env DJANGO_SETTINGS_MODULE=mysite.settings \
--master --pidfile=/tmp/project-master.pid \
--socket=127.0.0.1:49152 \      # can also be a file
--processes=5 \                  # number of worker processes
--uid=1000 --gid=2000 \          # if root, uwsgi can drop privileges
--harakiri=20 \                  # respawn processes taking more than 20 seconds
--max-requests=5000 \            # respawn processes after serving 5000 requests
--vacuum \                      # clear environment on exit
--home=/path/to/virtual/env \   # optional path to a virtualenv
--daemonize=/var/log/uwsgi/yourproject.log # background the process
```

**This is a lot to get right (or wrong)**

**Why is it so complicated?**

- Established tools like Gunicorn and uWSGI were designed before containers/Kubernetes
- A lot of stuff is not necessary for containers
  - Too many adjustable settings for beginners
  - Process management handled by Docker or Kubernetes
- Expecting a traditional web server like Nginx is no longer a fair assumption
  - We're not hosting LAMP servers
- Learning curve from dev server to production-ready is steep (Kubernetes or not)

# ASGI



- Asynchronous Server Gateway Interface
- async/await support: WebSockets & HTTP/2
- Early ASGI server implementations:
  - Daphne
  - Uvicorn
  - Hypercorn
- Needs new frameworks:
  - Django Channels, Django 3 (coming soon)
  - Quart (async Flask port)
  - Starlette, Sanic, ...
- Async ecosystem still evolving



 @jayhewland

 jamie.hewland@hpe.com

#### Further reading:

- *Deploying Django web applications in Docker containers*: [youtu.be/T2hooQzvurQ](https://youtu.be/T2hooQzvurQ)
- *Kubernetes, Local to Production with Django* by Mark Gituma: [bit.ly/k8slocaltoprod](https://bit.ly/k8slocaltoprod)

---

# Thank you

**aruba**  
a Hewlett Packard  
Enterprise company